



Time Series Analysis & Forecasting Using R

10. Forecast reconciliation



Instructors



Mitchell O'Hara-Wild

🏠 mitchelloharawild.com

⌚ [@mitchelloharawild](https://twitter.com/mitchelloharawild)

✉️ Mitch.OHara-Wild@monash.edu

Assumptions

- This is not an introduction to R. We assume you are broadly comfortable with R code, the RStudio environment and the tidyverse.
- This is not a statistics course. We assume you are familiar with concepts such as the mean, standard deviation, quantiles, regression, normal distribution, likelihood, etc.
- This is not a theory course. We are not going to derive anything. We will teach you time series and forecasting tools, when to use them, and how to use them most effectively.

Key reference

Hyndman, R. J. & Athanasopoulos, G. (2021)
***Forecasting: principles and practice*, 3rd ed.**

Key reference

Hyndman, R. J. & Athanasopoulos, G. (2021)
Forecasting: principles and practice, 3rd ed.

[OTexts.org/fpp3/](https://otexts.org/fpp3/)

Key reference

Hyndman, R. J. & Athanasopoulos, G. (2021)
***Forecasting: principles and practice*, 3rd ed.**

OTexts.org/fpp3/

- Free and online
- Data sets in associated R package
- R code for examples

Poll: How experienced are you in forecasting

- 1 Guru: I wrote the software, teach forecasting, and do the conference circuit.
- 2 Expert: It has been my full time job for more than a decade.
- 3 Skilled: I have been doing it for years.
- 4 Comfortable: I understand it and have done it.
- 5 Learner: I am still learning.
- 6 Beginner: I have heard of it and would like to learn more.
- 7 Unknown: What is forecasting? Is that what the weather people do?

Poll: How proficient are you in using R?

- 1 Guru: The R core team come to me for advice.
- 2 Expert: I have written several packages on CRAN.
- 3 Skilled: I use it regularly and it is an important part of my job.
- 4 Comfortable: I use it often and am comfortable with the tool.
- 5 User: I use it sometimes, but I am often searching around for the right function.
- 6 Learner: I have used it a few times.
- 7 Beginner: I've managed to download and install it.
- 8 Unknown: Why are you speaking like a pirate?

Install required packages

```
install.packages(c(  
  "tidyverse",  
  "fpp3",  
  "readabs",  
  "GGally",  
  "sugrrrants"  
))
```

Approximate outline

Day	Topic	Chapter
1	1. Introduction to tsibbles	2
1	2. Time series graphics	2
1	3. Transformations	3
1	4. Decomposing patterns	3
1	5. Introduction to forecasting	5,7
2	6. Exponential smoothing	8
2	7. ARIMA models	9,10
2	8. Accuracy evaluation	5
2	9. Multivariate modelling	12
2	10. Forecast reconciliation	11

workshop.nectric.com.au/tidyfc2024

Outline

- Time series data and tsibbles
 - Example: Australian prison population
 - Example: Australian pharmaceutical sales
 - Lab Session 1
 - Time plots
 - Lab Session 2
 - Seasonal plots
 - Lab Session 3
 - Seasonal or cyclic?
 - Lag plots and autocorrelation
 - Lab Session 4
 - White noise
 - Lab Session 5
 - Per capita adjustments
 - Lab Session 6
 - Inflation adjustments
 - Mathematical transformations
 - Lab Session 7
 - Time series decompositions
 - Lab Session 8
 - Multiple seasonality
 - A cautionary tale
 - Notice: Material planned to change
 - STL Features
 - Lab Session 9
 - Dimension reduction for features
 - Lab Session 10

Outline

- 1 Time series data and tsibbles
 - Example: Australian prison population
 - Example: Australian pharmaceutical sales
- 2 Lab Session 1
 - Time plots
- 3 Lab Session 2
 - Seasonal plots
- 4 Lab Session 3
 - Seasonal or cyclic?
- 5 Lag plots and autocorrelation
 - Lab Session 4
- 6 White noise
 - Lab Session 5
- 7 Per capita adjustments
 - Lab Session 6
- 8 Inflation adjustments
 - Mathematical transformations
- 9 Lab Session 7
 - Time series decompositions
- 10 Lab Session 8
 - Multiple seasonality
- 11 A cautionary tale
 - Notice: Material planned to change
- 12 STL Features
 - Lab Session 9
- 13 Dimension reduction for features
 - Lab Session 10

Tidyverts packages

tidyverts.org



Tidyverts developers

Earo Wang



Mitchell O'Hara-Wild



Time series data

- Four-yearly Olympic winning times
- Annual Google profits
- Quarterly Australian beer production
- Monthly rainfall
- Weekly retail sales
- Daily IBM stock prices
- Hourly electricity demand
- 5-minute freeway traffic counts
- Time-stamped stock transaction data

Class packages

```
# Data manipulation
library(dplyr)
# Plotting functions
library(ggplot2)
# Time and date manipulation
library(lubridate)
# Time series class
library(tsibble)
# Tidy time series data
library(tsibbldata)
# Time series graphics and statistics
library(feasts)
# Forecasting functions
library(fable)
```

Class packages

```
# Data manipulation
library(dplyr)
# Plotting functions
library(ggplot2)
# Time and date manipulation
library(lubridate)
# Time series class
library(tsibble)
# Tidy time series data
library(tsibbldata)
# Time series graphics and statistics
library(feasts)
# Forecasting functions
library(fable)
```

```
# All of the above
library(fpp3)
```

tsibble objects

```
{r, echo = TRUE} global_economy
```

tsibble objects

```
{r, echo = TRUE} global_economy
```

Index

tsibble objects

```
{r, echo = TRUE} global_economy
```

Index

Key

tsibble objects

```
{r, echo = TRUE} global_economy
```

Index

Key

Measured variables

tsibble objects

```
{r, echo = TRUE} tourism
```

Domestic visitor
nights in
thousands by
state/region and
purpose.

tsibble objects

```
{r, echo = TRUE} tourism
```

Index

Domestic visitor
nights in
thousands by
state/region and
purpose.

tsibble objects

```
{r, echo = TRUE} tourism
```

Index

Keys

Domestic visitor
nights in
thousands by
state/region and
purpose.

tsibble objects

```
{r, echo = TRUE} tourism
```

Index

Keys

Measure

Domestic visitor
nights in
thousands by
state/region and
purpose.

tsibble objects

- A tsibble allows storage and manipulation of multiple time series in R.
- It contains:
 - ▶ An index: time information about the observation
 - ▶ Measured variable(s): numbers of interest
 - ▶ Key variable(s): optional unique identifiers for each series
- It works with tidyverse functions.

The tsibble index

Example

```
mydata <- tsibble(  
  year = 2012:2016,  
  y = c(123, 39, 78, 52, 110),  
  index = year  
)  
mydata
```

The tsibble index

For observations more frequent than once per year, we need to use a time class function on the index.

```
{r tstablemonth, echo=FALSE} z <- tibble(Month = paste(2019,  
month.abb[1:5]), Observation = c(50, 23, 34, 30, 25)) #  
knitr::kable(z, booktabs=TRUE)
```

z

The tsibble index

For observations more frequent than once per year, we need to use a time class function on the index.

```
z |>  
  mutate(Month = yearmonth(Month)) |>  
  as_tsibble(index = Month)
```

The `tsibble` index

Common time index variables can be created with these functions:

```
{r tstable2, echo=FALSE} tribble(~`Frequency`,  
~Function, "Annual", "`start:end`", "Quarterly",  
"`yearquarter()`", "Monthly", "`yearmonth()`",  
"Weekly", "`yearweek()`", "Daily", "`as_date()`",  
`ymd()`", "Sub-daily", "`as_datetime()`" ) |>  
knitr::kable(booktabs = TRUE)
```

Outline

- Time series data and tsibbles
Example: Australian prison population
Example: Australian pharmaceutical sales
Lab Session 1
Time plots
Lab Session 2
Seasonal plots
Lab Session 3
Seasonal or cyclic?
Lag plots and autocorrelation
Lab Session 4
White noise
Lab Session 5
Per capita adjustments
Lab Session 6
Inflation adjustments
Mathematical transformations
Lab Session 7
Time series decompositions
Lab Session 8
Multiple seasonality
A cautionary tale
Notice: Material planned to change
STL Features
Lab Session 9
Dimension reduction for features
Lab Session 10

Australian prison population



Read a csv file and convert to a tsibble

```
prison <- readr::read_csv("data/prison_population.csv")
```

```
{r prison2a, dependson="prison", echo=FALSE} prison ## Read a csv file and convert  
to a tsibble
```

```
prison <- readr::read_csv("data/prison_population.csv") |>  
  mutate(Quarter = yearquarter(date))
```

```
{r prison3a, dependson="prison3", echo=FALSE} prison
```

Read a csv file and convert to a tsibble

```
prison <- readr::read_csv("data/prison_population.csv") |>  
  mutate(Quarter = yearquarter(date)) |>  
  select(-date)
```

```
{r prison4a, dependson="prison4", echo=FALSE} prison
```

Read a csv file and convert to a tsibble

```
prison <- readr::read_csv("data/prison_population.csv") |>  
  mutate(Quarter = yearquarter(date)) |>  
  select(-date) |>  
  as_tsibble(  
    index = Quarter,  
    key = c(state, gender, legal, indigenous)  
)
```

```
{r prison5a, dependson="prison5", echo=FALSE} prison
```

Outline

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

19

Time series data and tsibbles

Example: Australian prison population

Example: Australian pharmaceutical sales

Lab Session 1

Time plots

Lab Session 2

Seasonal plots

Lab Session 3

Seasonal or cyclic?

Lag plots and autocorrelation

Lab Session 4

White noise

Lab Session 5

Per capita adjustments

Lab Session 6

Inflation adjustments

Mathematical transformations

Lab Session 7

Time series decompositions

Lab Session 8

Multiple seasonality

A cautionary tale

Notice: Material planned to change

STL features

Lab Session 9

Dimension reduction for features

Australian Pharmaceutical Benefits Scheme



Australian Pharmaceutical Benefits Scheme

The **Pharmaceutical Benefits Scheme** (PBS) is the Australian government drugs subsidy scheme.

Australian Pharmaceutical Benefits Scheme

The **Pharmaceutical Benefits Scheme** (PBS) is the Australian government drugs subsidy scheme.

- Many drugs bought from pharmacies are subsidised to allow more equitable access to modern drugs.
- The cost to government is determined by the number and types of drugs purchased. Currently nearly 1% of GDP.
- The total cost is budgeted based on forecasts of drug usage.
- Costs are disaggregated by drug type (ATC1 \times length(unique(PBS\$ATC1)) / ATC2 \times length(unique(PBS\$ATC2))), concession category (x \times length(unique(PBS\$Concession))) and patient type (x \times length(unique(PBS\$Type))) giving $84 \times 2 \times 2 = 'r84 * 2 * 2'$

Working with tsibble objects

```
{r wide, include=FALSE} options(width = 78)
```

```
{r pbs1, dependson='wide'} PBS
```

Working with tsibble objects

We can use the `filter()` function to select rows.

```
PBS |>  
  filter(ATC2 == "A10")
```

Working with tsibble objects

We can use the `select()` function to select columns.

```
PBS |>
  filter(ATC2 == "A10") |>
  select(Month, Concession, Type, Cost)
```

Working with tsibble objects

We can use the `summarise()` function to summarise over keys.

```
PBS |>
  filter(ATC2 == "A10") |>
  select(Month, Concession, Type, Cost) |>
  summarise(total_cost = sum(Cost))
```

Working with tsibble objects

We can use the `mutate()` function to create new variables.

```
PBS |>
  filter(ATC2 == "A10") |>
  select(Month, Concession, Type, Cost) |>
  summarise(total_cost = sum(Cost)) |>
  mutate(total_cost = total_cost / 1e6)
```

Working with tsibble objects

We can use the `mutate()` function to create new variables.

```
PBS |>
  filter(ATC2 == "A10") |>
  select(Month, Concession, Type, Cost) |>
  summarise(total_cost = sum(Cost)) |>
  mutate(total_cost = total_cost / 1e6) -> a10
```

```
{r a10, echo=FALSE, dependson="pbs6"} a10
```

```
{r narrow, include=FALSE} options(width = 60)
```

Outline



Time series data and tsibbles

Example: Australian prison population

Example: Australian pharmaceutical sales

Lab Session 1

Time plots

Lab Session 2

Seasonal plots

Lab Session 3

Seasonal or cyclic?

Lag plots and autocorrelation

Lab Session 4

White noise

Lab Session 5

Per capita adjustments

Lab Session 6

Inflation adjustments

Mathematical transformations

Lab Session 7

Time series decompositions

Lab Session 8

Multiple seasonality

A cautionary tale

Notice: Material planned to change

STL features

Lab Session 9

Dimension reduction for features

Lab Session 1a

- 1 Use the `readabs` package to download and tidy the time series data from table "19" of the "6202.0" catalogue.
- 2 Use `filter()` to remove time series relating to hours worked disaggregated by State/Territory. Also remove the aggregated time series (which specify "Persons").
- 3 Split the series column into two columns, workload and sex using the `separate()` function.
- 4 Create a tsibble with the appropriate index and key variables, you will need to convert the date variable into a time variable with the appropriate granularity.

Lab Session 1b

Now let's look at this data:

- 5 Find which series contains the most worked hours.
- 6 Create a new tsibble which aggregates over sex, and just has monthly worked hours by type.
- 7 What is the typical difference in total working hours between full time and part time employees? (try to visualise this!)

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL Features
- Lab Session 9
- Dimension reduction for features

Time plots

```
{r maxtemp, fig.height=2.4} maxtemp <- vic_elec |> index_by(Day = date(Time)) |> summarise(Temperature = max(Temperature)) maxtemp |> autoplot(Temperature) + labs(y = "Max temperature")
```

Ansett airlines



Ansett airlines

```
{r, echo=TRUE, fig.height=3.1} ansett |> autoplot(Passengers)
```

Ansett airlines

```
{r, echo=TRUE, fig.height=3.1} ansett |> filter(Class == "Economy")  
|> autoplot(Passengers)
```

Ansett airlines

```
{r, echo=TRUE, fig.height=3.1} ansett |> filter(Airports ==  
"MEL-SYD") |> autoplot(Passengers)
```

Ansett airlines

```
{r, echo=TRUE, fig.height=3.1} ansett |> filter(Airports ==  
"MEL-SYD") |> autoplot(Passengers)
```

Not the real
data! Or is it?

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- 2 Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL Features
- Lab Session 9
- Dimension reduction for features

Lab Session 2

- Create time plots of the following four time series: Bricks from aus_production, Lynx from pelt, Close from gafa_stock, Demand from vic_elec.
- Use help() to find out about the data in each series.
- For the last plot, modify the axis labels and title.

Outline

- 1 Time series data and tsibbles
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Lab Session 1
- 5 Time plots
- 6 Lab Session 2
- 7 Seasonal plots
- 8 Lab Session 3
- 9 Seasonal or cyclic?
- 10 Lag plots and autocorrelation
- 11 Lab Session 4
- 12 White noise
- 13 Lab Session 5
- 14 Per capita adjustments
- 15 Lab Session 6
- 16 Inflation adjustments
- 17 Mathematical transformations
- 18 Lab Session 7
- 19 Time series decompositions
- 20 Lab Session 8
- 21 Multiple seasonality
- 22 A cautionary tale
- 23 Notice: Material planned to change
- 24 STL Features
- 25 Lab Session 9
- 26 Dimension reduction for features

Outline

- 1 Time series data and tsibbles
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Lab Session 1
- 5 Time plots
- 6 Lab Session 2
- 7 Seasonal plots
- 8 Lab Session 3
- 9 Seasonal or cyclic?
- 10 Lag plots and autocorrelation
- 11 Lab Session 4
- 12 White noise
- 13 Lab Session 5
- 14 Per capita adjustments
- 15 Lab Session 6
- 16 Inflation adjustments
- 17 Mathematical transformations
- 18 Lab Session 7
- 19 Time series decompositions
- 20 Lab Session 8
- 21 Multiple seasonality
- 22 A cautionary tale
- 23 Notice: Material planned to change
- 24 STL Features
- 25 Lab Session 9
- 26 Dimension reduction for features

Seasonal plots

- Data plotted against the individual “seasons” in which the data were observed. (In this case a “season” is a month.)
- Something like a time plot except that the data from each season are overlapped.
- Enables the underlying seasonal pattern to be seen more clearly, and also allows any substantial departures from the seasonal pattern to be easily identified.
- In R: `gg_season()`

Quarterly Australian Beer Production

```
{r, fig.height=2.7} beer <- aus_production |>  
select(Quarter, Beer) |> filter(year(Quarter) >=  
1992) beer |> autoplot(Beer)
```

Quarterly Australian Beer Production

```
beer |> gg_season(Beer, labels = "right")
```

Multiple seasonal periods

vic_elec

Multiple seasonal periods

```
{r, dev = "ragg_png", dpi = 180} vic_elec |>  
gg_season(Demand)
```

Multiple seasonal periods

```
{r, dev = "ragg_png", dpi = 180} vic_elec |>  
gg_season(Demand, period = "week")
```

Multiple seasonal periods

```
{r, dev = "ragg_png", dpi = 180} vic_elec |>  
gg_season(Demand, period = "day")
```

Seasonal subseries plots

- Data for each season collected together in time plot as separate time series.
- Enables the underlying seasonal pattern to be seen clearly, and changes in seasonality over time to be visualized.
- In R: `gg_subseries()`

Quarterly Australian Beer Production

```
beer |> gg_subseries(Beer)
```

Calendar plots

```
{r sugrrants, eval=FALSE} library(sugrrants) vic_elec  
|> filter(year(Date) == 2014) |> mutate(Hour =  
hour(Time)) |> frame_calendar(x = Hour, y = Demand,  
date = Date, nrow = 4) |> ggplot(aes(x = .Hour, y =  
.Demand, group = Date)) + geom_line() -> p1  
pretty(p1, size = 3, label.padding = unit(0.15,  
"lines")) )
```

- `frame_calendar()` makes a compact calendar plot
- `facet_calendar()` provides an easier `ggplot2` integration.

Calendar plots

```
{r sugrrants2, ref.label="sugrrants", echo=FALSE, out.height="90%",  
fig.height=5.4, fig.width=9}
```

Outline

- Time series data and tsibbles
Example: Australian prison population
Example: Australian pharmaceutical sales
Lab Session 1
Time plots
Lab Session 2
Seasonal plots
Lab Session 3
Seasonal or Cyclic?
Lag plots and autocorrelation
Lab Session 4
White noise
Lab Session 5
Per capita adjustments
Lab Session 6
Inflation adjustments
Mathematical transformations
Lab Session 7
Time series decompositions
Lab Session 8
Multiple seasonality
A cautionary tale
Notice: Material planned to change
STL Features
Lab Session 9
Dimension reduction for features
Lab Session 10

Lab Session 3

- 1 Look at the quarterly tourism data for the Snowy Mountains

```
snowy <- tourism |>  
  filter(Region == "Snowy Mountains")
```

- ▶ Use autoplot(), gg_season() and gg_subseries() to explore the data.
- ▶ What do you learn?

- 2 Produce a calendar plot for the pedestrian data from one location and one year.

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Time series patterns

Trend pattern exists when there is a long-term increase or decrease in the data.

Seasonal pattern exists when a series is influenced by seasonal factors (e.g., the quarter of the year, the month, or day of the week).

Cyclic pattern exists when data exhibit rises and falls that are *not of fixed period* (duration usually of at least 2 years).

Time series components

Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern

Time series patterns

```
{r, fig.height=2.7} aus_production |> filter(year(Quarter) >= 1980) |>  
autoplot(Electricity) + labs(y = "GWh", title = "Australian electricity  
production")
```

Time series patterns

```
{r, warning=FALSE, fig.height=2.7} aus_production |> autoplot(Bricks) +  
  labs(title = "Australian clay brick production", x = "Year", y =  
    "million units")
```

Time series patterns

```
{r, fig.height=2.7} us_employment |> filter>Title == "Retail Trade",  
year(Month) >= 1980) |> autoplot(Employed / 1e3) + labs(title = "Retail  
employment, USA", y = "Million people")
```

Time series patterns

```
{r, fig.height=2.7} gafa_stock |> filter(Symbol == "AMZN", year(Date) >= 2018) |> autoplot(Close) + labs(title = "Amazon closing stock price", x = "Day", y = "$")
```

Time series patterns

```
{r, fig.height=2.7} pelt |> autoplot(Lynx) + labs(title = "Annual Canadian  
Lynx Trappings", x = "Year", y = "Number trapped")
```

Seasonal or cyclic?

Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern

Seasonal or cyclic?

Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern

The timing of peaks and troughs is predictable with seasonal data, but unpredictable in the long term with cyclic data.

Outline

- 1 Time series data and tsibbles
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Lab Session 1
- 5 Time plots
- 6 Lab Session 2
- 7 Seasonal plots
- 8 Lab Session 3
- 9 Seasonal or cyclic?
- 10 Lag plots and autocorrelation
- 11 Lab Session 4
- 12 White noise
- 13 Lab Session 5
- 14 Per capita adjustments
- 15 Lab Session 6
- 16 Inflation adjustments
- 17 Mathematical transformations
- 18 Lab Session 7
- 19 Time series decompositions
- 20 Lab Session 8
- 21 Multiple seasonality
- 22 A cautionary tale
- 23 Notice: Material planned to change
- 24 STL Features
- 25 Lab Session 9
- 26 Dimension reduction for features

Example: Beer production

```
new_production <- aus_production |>  
  filter(year(Quarter) >= 1992)  
new_production
```

Example: Beer production

```
{r, fig.height=6.5, fig.width=6.5, out.width="7cm"}  
new_production |> gg_lag(Beer)
```

Example: Beer production

```
{r, fig.height=6.5, fig.width=6.5, out.width="7cm"}  
new_production |> gg_lag(Beer, geom = "point")
```

Lagged scatterplots

- Each graph shows y_t plotted against y_{t-k} for different values of k .
- The autocorrelations are the correlations associated with these scatterplots.
- ACF (autocorrelation function):
 - ▶ $r_1 = \text{Correlation}(y_t, y_{t-1})$
 - ▶ $r_2 = \text{Correlation}(y_t, y_{t-2})$
 - ▶ $r_3 = \text{Correlation}(y_t, y_{t-3})$
 - ▶ etc.
- If there is **seasonality**, the ACF at the seasonal lag (e.g., 12 for monthly data) will be **large and positive**.

Autocorrelation

Results for first 9 lags for beer data:

```
{r, echo=TRUE} new_production |> ACF(Beer, lag_max = 9)
```

Autocorrelation

Results for first 9 lags for beer data:

```
{r beeracf, fig.height=2.5} new_production |> ACF(Beer, lag_max = 9) |>  
autoplot()
```

ACF

```
{r, fig.height=3, echo=TRUE} new_production |>  
ACF(Beer) |> autoplot()
```

Australian student enrolments

```
students <- readxl::read_excel("../data/schools/Table 42b Number of Full-time and Pa  
group_by(Year, State = substring(`State/Territory`, 3)) |>  
summarise(Count = sum(`All Full-time and Part-time Student count`), .groups = "dro  
as_tsibble(index = Year, key = State)  
  
{r, echo=FALSE} students
```

Australian student enrolments

```
{r holidays-plot, echo=TRUE, dependson="holidays", fig.height=3.1} students |>  
autoplot(Count) + labs(y = "Student Count", title = "Australian students") +  
scale_y_log10()
```

Australian holidays

students |> ACF(Count)

Australian holidays

```
{r tourismacf2, fig.height=5, fig.width=5,  
out.width="49%"} students |> ACF(Count) |> autoplot()
```

Trend and seasonality in ACF plots

- When data have a trend, the autocorrelations for small lags tend to be large and positive.
- When data are seasonal, the autocorrelations will be larger at the seasonal lags (i.e., at multiples of the seasonal frequency)
- When data are trended and seasonal, you see a combination of these effects.

US retail trade employment

```
retail <- us_employment |>  
  filter>Title == "Retail Trade", year(Month) >= 1980)  
retail |> autoplot(Employed)
```

US retail trade employment

```
retail |>  
  ACF(Employed, lag_max = 48) |>  
  autoplot()
```

Google stock price

```
google_2015 <- gafa_stock |>  
  filter(Symbol == "GOOG", year(Date) == 2015) |>  
  select(Date, Close)  
google_2015
```

Google stock price

```
google_2015 |> autoplot(Close)
```

Google stock price

```
google_2015 |>  
  ACF(Close, lag_max = 100) |>  
  autoplot()
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Lab Session 4

We have introduced the following functions: `gg_lag` and `ACF`. Use these functions to explore the four time series: Bricks from `aus_production`, Lynx from `pelt`, Close price of Amazon from `gafa_stock`, Demand from `vic_elec`. Can you spot any seasonality, cyclicity and trend? What do you learn about the series?

Which is which?

```
{r, fig.height=6, fig.width=12, echo=FALSE,  
warning=FALSE, out.width="100%"} cowtemp <-  
as_tsibble(fma::cowtemp) USAccDeaths <-  
as_tsibble(USAccDeaths) AirPassengers <-  
as_tsibble(AirPassengers) mink <-  
as_tsibble(fma::mink) tp1 <- autoplot(cowtemp, value)  
+ labs(x = "") + labs(y = "chirps per minute") +  
labs(title = "1. Daily temperature of cow") tp2 <-  
autoplot(USAccDeaths, value) + labs(x = "") + labs(y =  
"thousands") + labs(title = "2. Monthly accidental  
deaths") tp3 <- autoplot(AirPassengers, value) +
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Faq plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL Features
- Lab Session 9
- Dimension reduction for features

Example: White noise

```
{r, fig.height=3.3} wn <- tsibble(t = seq(36), y = rnorm(36), index = t) wn |> autoplot(y)
```

Example: White noise

```
{r, fig.height=3.3} wn <- tsibble(t = seq(36), y = rnorm(36), index = t) wn |> autoplot(y)
```

White noise data is uncorrelated across time with zero mean and constant variance.
(Technically, we require independence as well.)

Example: White noise

```
wn |> ACF(y)
```

```
{r wnacf, echo=FALSE, dependson="wn"} wn |> ACF(y, lag_max = 10) |>  
as_tibble() |> mutate(lag = as.numeric(lag)) |> pivot_wider(names_from =  
lag, values_from = acf) |> rename_all(function(x) { paste("$r_{", x,  
"}$", sep = "") }) |> knitr::kable( booktabs = TRUE, escape =  
FALSE, align = "c", digits = 3, format.args = list(nsmall = 3) )  
  
{r, echo=FALSE, fig.height=1.5} wn |> ACF(y) |> autoplot()
```

Example: White noise

```
wn |> ACF(y)
```

```
{r wnacf, echo=FALSE, dependson="wn"} wn |> ACF(y, lag_max = 10) |>  
as_tibble() |> mutate(lag = as.numeric(lag)) |> pivot_wider(names_from =  
lag, values_from = acf) |> rename_all(function(x) { paste("$r_{", x,  
"}$"), sep = "") }) |> knitr::kable( booktabs = TRUE, escape =  
FALSE, align = "c", digits = 3, format.args = list(nsmall = 3) )  
  
{r, echo=FALSE, fig.height=1.5} wn |> ACF(y) |> autoplot()
```

- Sample autocorrelations for white noise series.
- Expect each autocorrelation to be close to zero.
- Blue lines show 95% critical values.

Example: Pigs slaughtered

```
{r, fig.height=2.5} pigs <- aus_livestock |> filter(State ==  
"Victoria", Animal == "Pigs", year(Month) >= 2014) pigs |>  
autoplot(Count / 1e3) + labs(x = "Year", y = "Thousands",  
title = "Number of pigs slaughtered in Victoria")
```

Example: Pigs slaughtered

```
pigs |>  
  ACF(Count) |>  
  autoplot()
```

Example: Pigs slaughtered

Monthly total number of pigs slaughtered in the state of Victoria, Australia, from January 2014 through December 2018
(Source: Australian Bureau of Statistics.)

Example: Pigs slaughtered

Monthly total number of pigs slaughtered in the state of Victoria, Australia, from January 2014 through December 2018
(Source: Australian Bureau of Statistics.)

- Difficult to detect pattern in time plot.
- ACF shows significant autocorrelation for lag 2 and 12.
- Indicate some slight seasonality.

Example: Pigs slaughtered

Monthly total number of pigs slaughtered in the state of Victoria, Australia, from January 2014 through December 2018
(Source: Australian Bureau of Statistics.)

- Difficult to detect pattern in time plot.
- ACF shows significant autocorrelation for lag 2 and 12.
- Indicate some slight seasonality.

These show the series is **not a white noise series**.

Outline

- 1 Time series data and tsibbles
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Lab Session 1
- 5 Time plots
- 6 Lab Session 2
- 7 Seasonal plots
- 8 Lab Session 3
- 9 Seasonal or cyclic?
- 10 Lag plots and autocorrelation
- 11 Lab Session 4
- 12 White noise
- 13 **Lab Session 5**
- 14 Per capita adjustments
- 15 Lab Session 6
- 16 Inflation adjustments
- 17 Mathematical transformations
- 18 Lab Session 7
- 19 Time series decompositions
- 20 Lab Session 8
- 21 Multiple seasonality
- 22 A cautionary tale
- 23 Notice: Material planned to change
- 24 STL features
- 25 Lab Session 9
- 26 Dimension reduction for features

Lab Session 5

Calculate the difference in ACT student enrolments, it can be done as follows:

```
{r, eval = FALSE} students |> filter(State == "ACT") |> mutate(diff = difference(Count))
```

Does diff look like white noise?

Outline

- 1 Time series data and tsibbles
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Lab Session 1
- 5 Time plots
- 6 Lab Session 2
- 7 Seasonal plots
- 8 Lab Session 3
- 9 Seasonal or cyclic?
- 10 Lag plots and autocorrelation
- 11 Lab Session 4
- 12 White noise
- 13 Lab Session 5
- 14 Per capita adjustments
- 15 Lab Session 6
- 16 Inflation adjustments
- 17 Mathematical transformations
- 18 Lab Session 7
- 19 Time series decompositions
- 20 Lab Session 8
- 21 Multiple seasonality
- 22 A cautionary tale
- 23 Notice: Material planned to change
- 24 STL Features
- 25 Lab Session 9
- 26 Dimension reduction for features

Outline

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

19

Time series data and tsibbles

Example: Australian prison population

Example: Australian pharmaceutical sales

Lab Session 1

Time plots

Lab Session 2

Seasonal plots

Lab Session 3

Seasonal or cyclic?

Lag plots and autocorrelation

Lab Session 4

White noise

Lab Session 5

Per capita adjustments

Lab Session 6

Inflation adjustments

Mathematical transformations

Lab Session 7

Time series decompositions

Lab Session 8

Multiple seasonality

A cautionary tale

Notice: Material planned to change

STL features

Lab Session 9

Dimension reduction for features

Per capita adjustments

```
global_economy |>  
  filter(Country == "Australia") |>  
  autoplot(GDP)
```

Per capita adjustments

```
global_economy |>  
  filter(Country == "Australia") |>  
  autoplot(GDP / Population)
```

Outline

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

19

Time series data and tsibbles

Example: Australian prison population

Example: Australian pharmaceutical sales

Lab Session 1

Time plots

Lab Session 2

Seasonal plots

Lab Session 3

Seasonal or cyclic?

Lag plots and autocorrelation

Lab Session 4

White noise

Lab Session 5

Per capita adjustments

Lab Session 6

Inflation adjustments

Mathematical transformations

Lab Session 7

Time series decompositions

Lab Session 8

Multiple seasonality

A cautionary tale

Notice: Material planned to change

STL features

Lab Session 9

Dimension reduction for features

Lab Session 6

Consider the GDP information in `global_economy`. Plot the GDP per capita for each country over time. Which country has the highest GDP per capita? How has this changed over time?

Outline



Time series data and tsibbles
Example: Australian prison population
Example: Australian pharmaceutical sales
Lab Session 1
Time plots
Lab Session 2
Seasonal plots
Lab Session 3
Seasonal or cyclic?
Lag plots and autocorrelation
Lab Session 4
White noise
Lab Session 5
Per capita adjustments
Lab Session 6
Inflation adjustments
Mathematical transformations
Lab Session 7
Time series decompositions
Lab Session 8
Multiple seasonality
A cautionary tale
Notice: Material planned to change
STL Features
Lab Session 9
Dimension reduction for features

Inflation adjustments

```
{r, eval=FALSE} print_retail <- aus_retail |> filter(Industry == "Newspaper and book retailing") |> group_by(Industry) |> index_by(Year = year(Month)) |> summarise(Turnover = sum(Turnover)) aus_economy <- filter(global_economy, Country == "Australia") print_retail |> left_join(aus_economy, by = "Year") |> mutate(Adj_turnover = Turnover / CPI) |> pivot_longer(c(Turnover, Adj_turnover), names_to = "Type", values_to = "Turnover" ) |> ggplot(aes(x = Year, y = Turnover)) + geom_line() + facet_grid(vars(Type), scales = "free_y") + labs(x = "Years", y = NULL, title = "Turnover: Australian print media industry")
```

Inflation adjustments

```
{r, message=FALSE, warning=FALSE, echo=FALSE, fig.height=5, out.height="90%"}

print_retail <- aus_retail |> filter(Industry == "Newspaper and book
retailing") |> group_by(Industry) |> index_by(Year = year(Month)) |>
summarise(Turnover = sum(Turnover)) aus_economy <-
filter(tsibbledata::global_economy, Country == "Australia") print_retail |>
left_join(aus_economy, by = "Year") |> mutate(Adj_turnover = Turnover / CPI)
|> pivot_longer(c(Turnover, Adj_turnover), names_to = "Type", values_to
= "Turnover" ) |> ggplot(aes(x = Year, y = Turnover)) + geom_line() +
facet_grid(vars(Type), scales = "free_y") + labs( x = "Years", y = NULL,
title = "Turnover: Australian print media industry" )
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Variance stabilization

If the data show different variation at different levels of the series, then a transformation can be useful.

Variance stabilization

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as y_1, \dots, y_n and transformed observations as w_1, \dots, w_n .

Variance stabilization

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as y_1, \dots, y_n and transformed observations as w_1, \dots, w_n .

Mathematical transformations for stabilizing variation

Square root $w_t = \sqrt{y_t}$ ↓

Cube root $w_t = \sqrt[3]{y_t}$ Increasing

Logarithm $w_t = \log(y_t)$ strength

Variance stabilization

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as y_1, \dots, y_n and transformed observations as w_1, \dots, w_n .

Mathematical transformations for stabilizing variation

Square root $w_t = \sqrt{y_t}$ ↓

Cube root $w_t = \sqrt[3]{y_t}$ Increasing

Logarithm $w_t = \log(y_t)$ strength

Logarithms, in particular, are useful because they are more interpretable: changes in a log value are **relative (percent) changes**

Variance stabilization

```
{r food, echo=TRUE} food <- aus_retail |> filter(Industry ==  
"Food retailing") |> summarise(Turnover = sum(Turnover))  
  
{r food-plot, echo = FALSE, fig.height=3.8} food |>  
autoplot(Turnover) + labs(y = "Turnover ($AUD)")
```

Variance stabilization

```
{r food-sqrt1, echo=TRUE, fig.height=3.8} food |>  
  autoplot(sqrt(Turnover)) +  labs(y = "Square root turnover")
```

Variance stabilization

```
{r food-cbrt, echo=TRUE, fig.height=3.8} food |>  
  autoplot(Turnover^(1 / 3)) +  labs(y = "Cube root turnover")
```

Variance stabilization

```
{r food-log, echo=TRUE, fig.height=3.8} food |>  
  autoplot(log(Turnover)) +  labs(y = "Log turnover")
```

Variance stabilization

```
{r food-inverse, echo=TRUE, fig.height=3.8} food |> autoplot(-1 /  
Turnover) + labs(y = "Inverse turnover")
```

Box-Cox transformations

Each of these transformations is close to a member of the family of
Box-Cox transformations:

$$w_t = \begin{cases} \log(y_t), & \lambda = 0; \\ (\text{sign}(y_t)|y_t|^\lambda - 1)/\lambda, & \lambda \neq 0. \end{cases}$$

Box-Cox transformations

Each of these transformations is close to a member of the family of **Box-Cox transformations**:

$$w_t = \begin{cases} \log(y_t), & \lambda = 0; \\ (\text{sign}(y_t)|y_t|^\lambda - 1)/\lambda, & \lambda \neq 0. \end{cases}$$

- Actually the Bickel-Doksum transformation (allowing for $y_t < 0$)
- $\lambda = 1$: (No substantive transformation)
- $\lambda = \frac{1}{2}$: (Square root plus linear transformation)
- $\lambda = 0$: (Natural logarithm)
- $\lambda = -1$: (Inverse plus 1)

Box-Cox transformations

```
{r food-anim, echo=FALSE, fig.show='animate',
interval=1/10, message=FALSE, fig.height=5,
fig.width=8,
aniopts='controls,buttonsize=0.3cm,width=11.5cm'}
library(rlang) library(gganimate) library(latex2exp)
food |>   mutate(!!!set_names(map(seq(0, 1, 0.01), ~
expr(fabletools::box_cox(Turnover, !!!x))), seq(0, 1,
0.01))) |>   select(-Turnover) |>
pivot_longer(-Month, names_to = "lambda", values_to =
"Turnover") |>   mutate(lambda = as.numeric(lambda))
|>   ggplot(aes(x = Month, y = Turnover)) +
```

Box-Cox transformations

```
{r food-lambda, echo=TRUE} food |> features(Turnover,  
features = guerrero)
```

Box-Cox transformations

```
{r food-lambda, echo=TRUE} food |> features(Turnover,  
features = guerrero)
```

- This attempts to balance the seasonal fluctuations and random variation across the series.
- Always check the results.
- A low value of λ can give extremely large prediction intervals.

Box-Cox transformations

```
{r food-bc, echo=TRUE,fig.height=3.8} food |>  
autoplot(box_cox(Turnover, 0.0895)) + labs(y = "Box-Cox  
transformed turnover")
```

Transformations

- Often no transformation needed.
- Simple transformations are easier to explain and work well enough.
- Transformations can have very large effect on PI.
- If some data are zero or negative, then use $\lambda > 0$.
- `log1p()` can also be useful for data with zeros.
- Choosing logs is a simple way to force forecasts to be positive
- Transformations must be reversed to obtain forecasts on the original scale. (Handled automatically by `fable`.)

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7**
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Lab Session 7

1 For the following series, find an appropriate transformation in order to stabilise the variance.

- ▶ United States GDP from `global_economy`
- ▶ Slaughter of Victorian “Bulls, bullocks and steers” in `aus_livestock`
- ▶ Victorian Electricity Demand from `vic_elec`.
- ▶ Gas production from `aus_production`

2 Why is a Box-Cox transformation unhelpful for the `canadian_gas` data?

Outline

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
289
290
291
292
293
294
295
296
297
298
299
299
300
301
302
303
304
305
306
307
308
309
309
310
311
312
313
314
315
316
317
318
319
319
320
321
322
323
324
325
326
327
328
329
329
330
331
332
333
334
335
336
337
338
339
339
340
341
342
343
344
345
346
347
348
349
349
350
351
352
353
354
355
356
357
358
359
359
360
361
362
363
364
365
366
367
368
369
369
370
371
372
373
374
375
376
377
378
379
379
380
381
382
383
384
385
386
387
388
389
389
390
391
392
393
394
395
396
397
398
399
399
400
401
402
403
404
405
406
407
408
409
409
410
411
412
413
414
415
416
417
418
419
419
420
421
422
423
424
425
426
427
428
429
429
430
431
432
433
434
435
436
437
438
439
439
440
441
442
443
444
445
446
447
448
449
449
450
451
452
453
454
455
456
457
458
459
459
460
461
462
463
464
465
466
467
468
469
469
470
471
472
473
474
475
476
477
478
479
479
480
481
482
483
484
485
486
487
488
489
489
490
491
492
493
494
495
496
497
498
499
499
500
501
502
503
504
505
506
507
508
509
509
510
511
512
513
514
515
516
517
518
519
519
520
521
522
523
524
525
526
527
528
529
529
530
531
532
533
534
535
536
537
538
539
539
540
541
542
543
544
545
546
547
548
549
549
550
551
552
553
554
555
556
557
558
559
559
560
561
562
563
564
565
566
567
568
569
569
570
571
572
573
574
575
576
577
578
579
579
580
581
582
583
584
585
586
587
588
589
589
590
591
592
593
594
595
596
597
598
598
599
599
600
601
602
603
604
605
606
607
608
609
609
610
611
612
613
614
615
616
617
618
619
619
620
621
622
623
624
625
626
627
628
629
629
630
631
632
633
634
635
636
637
638
639
639
640
641
642
643
644
645
646
647
648
649
649
650
651
652
653
654
655
656
657
658
659
659
660
661
662
663
664
665
666
667
668
669
669
670
671
672
673
674
675
676
677
678
679
679
680
681
682
683
684
685
686
687
688
689
689
690
691
692
693
694
695
696
697
697
698
699
699
700
701
702
703
704
705
706
707
708
709
709
710
711
712
713
714
715
716
717
718
719
719
720
721
722
723
724
725
726
727
728
729
729
730
731
732
733
734
735
736
737
738
739
739
740
741
742
743
744
745
746
747
748
749
749
750
751
752
753
754
755
756
757
758
759
759
760
761
762
763
764
765
766
767
768
769
769
770
771
772
773
774
775
776
777
778
779
779
780
781
782
783
784
785
786
787
788
789
789
790
791
792
793
794
795
796
797
797
798
799
799
800
801
802
803
804
805
806
807
808
809
809
810
811
812
813
814
815
816
817
818
819
819
820
821
822
823
824
825
826
827
828
829
829
830
831
832
833
834
835
836
837
838
839
839
840
841
842
843
844
845
846
847
848
849
849
850
851
852
853
854
855
856
857
858
859
859
860
861
862
863
864
865
866
867
868
869
869
870
871
872
873
874
875
876
877
878
879
879
880
881
882
883
884
885
886
887
888
889
889
890
891
892
893
894
895
896
897
897
898
899
899
900
901
902
903
904
905
906
907
908
909
909
910
911
912
913
914
915
916
917
918
919
919
920
921
922
923
924
925
926
927
928
929
929
930
931
932
933
934
935
936
937
938
939
939
940
941
942
943
944
945
946
947
948
949
949
950
951
952
953
954
955
956
957
958
959
959
960
961
962
963
964
965
966
967
968
969
969
970
971
972
973
974
975
976
977
978
979
979
980
981
982
983
984
985
986
987
988
989
989
990
991
992
993
994
995
995
996
997
997
998
999
999
1000

- Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Outline

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

19

Time series data and tsibbles

Example: Australian prison population

Example: Australian pharmaceutical sales

Lab Session 1

Time plots

Lab Session 2

Seasonal plots

Lab Session 3

Seasonal or cyclic?

Lag plots and autocorrelation

Lab Session 4

White noise

Lab Session 5

Per capita adjustments

Lab Session 6

Inflation adjustments

Mathematical transformations

Lab Session 7

Time series decompositions

Lab Session 8

Multiple seasonality

A cautionary tale

Notice: Material planned to change

STL features

Lab Session 9

Dimension reduction for features

Time series decomposition

Trend-Cycle aperiodic changes in level over time.

Seasonal (almost) periodic changes in level due to seasonal factors (e.g., the quarter of the year, the month, or day of the week).

Additive decomposition

$$y_t = S_t + T_t + R_t$$

where y_t = data at period t

T_t = trend-cycle component at period t

S_t = seasonal component at period t

R_t = remainder component at period t

STL decomposition

- STL: “Seasonal and Trend decomposition using Loess”
- Very versatile and robust.
- Seasonal component allowed to change over time, and rate of change controlled by user.
- Smoothness of trend-cycle also controlled by user.
- Optionally robust to outliers
- No trading day or calendar adjustments.
- Only additive.
- Take logs to get multiplicative decomposition.
- Use Box-Cox transformations to get other decompositions.

US Retail Employment

```
us_retail_employment <- us_employment |>
  filter(year(Month) >= 1990, Title == "Retail Trade") |>
  select(-Series_ID)
us_retail_employment
```

US Retail Employment

```
us_retail_employment |>  
  autoplot(Employed) +  
  labs(y = "Persons (thousands)", title = "Total employment in US retail")
```

US Retail Employment

```
dcmp <- us_retail_employment |>  
  model(stl = STL(Employed))  
dcmp
```

US Retail Employment

components (dcmp)

US Retail Employment

```
{r usretail-stl, fig.width=8, fig.height=4}  
components(dcmp) |> autoplot()
```

US Retail Employment

```
{r dable4, fig.height=2.7} us_retail_employment |>  
autoplot(Employed, color = "gray") +  
autolayer(components(dcmp), trend, color = "#D55E00")  
+ labs(y = "Persons (thousands)", title = "Total  
employment in US retail")
```

US Retail Employment

```
components(dcmp) |> gg_subseries(season_year)
```

Seasonal adjustment

- Useful by-product of decomposition: an easy way to calculate seasonally adjusted data.
- Additive decomposition: seasonally adjusted data given by

$$y_t - S_t = T_t + R_t$$

- Multiplicative decomposition: seasonally adjusted data given by

$$y_t/S_t = T_t \times R_t$$

US Retail Employment

```
{r usretail-sa, fig.height=2.7} us_retail_employment  
|> autoplot(Employed, color = "gray") +  
autolayer(components(dcmp), season_adjust, color =  
"#0072B2") + labs(y = "Persons (thousands)", title =  
"Total employment in US retail")
```

Seasonal adjustment

- We use estimates of S based on past values to seasonally adjust a current value.
- Seasonally adjusted series reflect **remainders** as well as **trend**. Therefore they are not “smooth” and “downturns” or “upturns” can be misleading.
- It is better to use the trend-cycle component to look for turning points.

STL decomposition

```
“{r stlwindowanim, echo=FALSE, warning=FALSE,
message=FALSE, fig.show='animate', interval=1/10,
fig.height=5.35, fig.width=8,
aniopts='controls,buttonsize=0.3cm,width=11.5cm',
eval=TRUE} s_windows <- seq(5, 55, by = 2) stl_defs <-
purrr::map(s_windows, function(s_window) { STL(Employed ~
season(window = s_window), robust = TRUE) }) names(stl_defs)
<- sprintf("season(window=%02d)", s_windows)

us_retail_employment |> model (!!stl_defs) |> components() |>
as_tibble() |> pivot_longer(Employed:remainder, names_to =
"component", names_ptypes = list(component = factor(levels =
```

STL decomposition

```
{r mstl, fig.width=8.5, fig.height=3.4}  
us_retail_employment |> model(STL(Employed)) |>  
components() |> autoplot()
```

STL decomposition

- STL() chooses season(window=13) by default
- Can include transformations.

```
{r mstl, fig.width=8.5, fig  
us_retail_employment |>    model(STL(Employed)) |>  
components() |>    autoplot()
```

STL decomposition

- Algorithm that updates trend and seasonal components iteratively.
- Starts with $\hat{T}_t = 0$
- Uses a mixture of loess and moving averages to successively refine the trend and seasonal estimates.
- trend window controls loess bandwidth on deasonalised values.
- season window controls loess bandwidth on detrended subseries.
- Robustness weights based on remainder.
- Default season: window = 13
- Default trend:

window =

```
nextodd(ceiling((1.5*period)/(1-(1.5/s.window))))
```

Australian holidays

```
{r holidays, include=FALSE} holidays <- tourism |>  
filter(Purpose == "Holiday") |> group_by(State) |>  
summarise(Trips = sum(Trips))  
  
{r holidays-plot2, echo=TRUE, dependson="holidays",  
fig.height=2.8} holidays |> autoplot(Trips) + labs(y  
= "thousands of trips", x = "Year", title =  
"Australian domestic holiday nights")
```

Australian holidays

```
{r stlagain2, echo=TRUE, warning=FALSE, fig.width=8,  
fig.height=3} holidays |> model(stl = STL(Trips)) |>  
components() |> autoplot()
```

Holidays decomposition

```
dcmp <- holidays |>  
  model(stl = STL(Trips)) |>  
  components()  
dcmp
```

Holidays decomposition

```
{r holidays3, fig.height=4.6} dcmp |> gg_subseries(season_year)
```

Holidays decomposition

```
{r holidays-trend, message=FALSE, warning=FALSE, fig.height=3.5, out.height="70%"}  
autoplot(dcmp, trend, scale_bars = FALSE) + autolayer(holidays, alpha = 0.4)
```

Outline



Time series data and tsibbles

Example: Australian prison population

Example: Australian pharmaceutical sales

Lab Session 1

Time plots

Lab Session 2

Seasonal plots

Lab Session 3

Seasonal or cyclic?

Lag plots and autocorrelation

Lab Session 4

White noise

Lab Session 5

Per capita adjustments

Lab Session 6

Inflation adjustments

Mathematical transformations

Lab Session 7

Time series decompositions

Lab Session 8

Multiple seasonality

A cautionary tale

Notice: Material planned to change

STL features

Lab Session 9

Dimension reduction for features

Lab Session 8

- 1 Produce the following decomposition

```
canadian_gas |>  
  model(STL(Volume ~ season(window=7) + trend(window=11))) |>  
  components() |>  
  autoplot()
```

- 2 What happens as you change the values of the two window arguments?
- 3 How does the seasonal shape change over time? [Hint: Try plotting the seasonal component using gg_season.]
- 4 Can you produce a plausible seasonally adjusted series? [Hint: season_adjust is one of the variables returned by STL.]

Outline

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

19

20

21

22

Time series data and tsibbles

Example: Australian prison population

Example: Australian pharmaceutical sales

Lab Session 1

Time plots

Lab Session 2

Seasonal plots

Lab Session 3

Seasonal or cyclic?

Lag plots and autocorrelation

Lab Session 4

White noise

Lab Session 5

Per capita adjustments

Lab Session 6

Inflation adjustments

Mathematical transformations

Lab Session 7

Time series decompositions

Lab Session 8

Multiple seasonality

A cautionary tale

Notice: Material planned to change

STL features

Lab Session 9

Dimension reduction for features

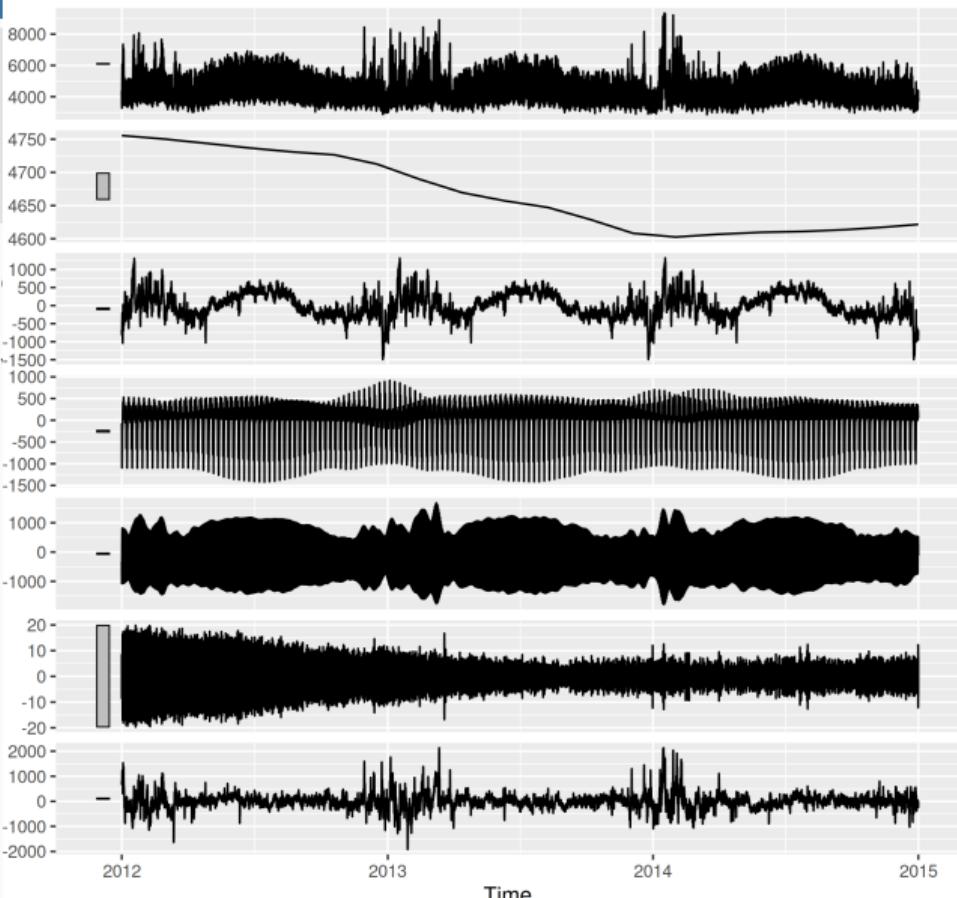
Multiple seasonality

```
vic_elec |>  
  model(STL(Demand)) |>  
  components() |>  
  autoplot()
```

```
{r vic_elec, echo=FALSE} png("figs/vi  
300, units = "cm", type = "cairo-png"  
components() |>  autoplot() crop::de
```

STL decomposition

Demand = trend + season_year + season_week + season_day + season_hour + remainder



Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL Features
- Lab Session 9
- Dimension reduction for features

A cautionary tale

NEWS 

LOCATION: Clayton, Vic [Change](#)

[Home](#) Just In Australia World Business Sport Analysis & Opinion Fact Check Programs

BREAKING NEWS

Police arrest man in connection with stabbing death of 17-year-old Masa Vukotic in M

 Print  Email  Facebook  Twitter  More

Treasurer Joe Hockey calls for answers over Australian Bureau of Statistics jobs data

By [Michael Vincent](#) and [Simon Frazer](#)

Updated 9 Oct 2014, 12:17pm

Federal Treasurer Joe Hockey says he wants answers to the problems the Australian Bureau of Statistics (ABS) has had with unemployment figures.

Mr Hockey, who is in the US to discuss Australia's G20 agenda, said last month's unemployment figures were "extraordinary".

The rate was 6.1 per cent after jumping to a 12-year high of 6.4 per cent the previous month.

The ABS has now taken the rare step of abandoning seasonal adjustment for its latest employment data.



PHOTO: Joe Hockey says he is unhappy with the volatility of ABS unemployment figures. (AAP: Alan Porritt)

RELATED STORY: ABS abandons seasonal adjustment for latest jobs data

A cautionary tale

NEWS 

LOCATION: Clayton, Vic [Change](#)

 [ABC TV](#) [ABC Radio](#)

[Home](#) [Just In](#) [Australia](#) [World](#) [Business](#) [Sport](#) [Analysis & Opinion](#) [Fact Check](#) [Programs](#) [Media](#) [Topics](#) [Search](#)

BREAKING NEWS Police arrest man in connection with stabbing death of 17-year-old Masa Vukotic in Melb

[!\[\]\(6b4a37805d7af32df4ed90f90d37d246_img.jpg\) Print](#) [!\[\]\(a4281bd267652205907243eee7c0fc94_img.jpg\) Email](#) [!\[\]\(37ae4c9e77aac80be07abdbc5997f413_img.jpg\) Facebook](#) [!\[\]\(0e7a00a5d4390c826d84f8849a4e6e37_img.jpg\) Twitter](#) [!\[\]\(038d31f9fca9f706039b78979d067859_img.jpg\) More](#)

ABS abandons seasonal adjustment for latest jobs data

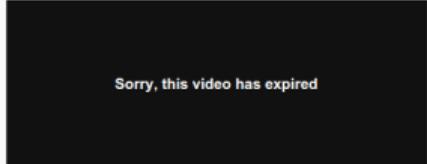
By business reporter [Michael Janda](#)
Updated 8 Oct 2014, 4:19pm

The Australian Bureau of Statistics is taking the rare step of abandoning seasonal adjustment for its latest employment data.

The ABS uses seasonal adjustment, based on historical experience, to account for the normal variation between hiring and firing patterns between different months.

However, after a winter where the seasonally adjusted unemployment rate swung wildly from 6.1 to 6.4 and back to 6.1 per cent, [the bureau released a statement](#) saying it will not adjust the original figure for September for seasonal factors.

It will also reset the seasonal adjustment for July and August to one, meaning that these months will also reflect the original figures.



Sorry, this video has expired

VIDEO: Westpac chief economist Bill Evans discusses the ABS jobs data changes (ABC News)

RELATED STORY: Doubt the record breaking jobs figures? So does the ABS

RELATED STORY: Jobs increase record sees unemployment slashed

RELATED STORY: Unemployment surges to 12-year high at 6.4 pc

MAP: Australia 

A cautionary tale

ABS jobs and unemployment figures - key questions answered by an expert

A professor of statistics at Monash University explains exactly what is seasonal adjustment, why it matters and what went wrong in the July and August figures



© School leavers come on to the jobs market at the same time, causing a seasonal fluctuation. Photograph: Brian Snyder/Reuters

The Australian Bureau of Statistics has retracted its seasonally adjusted employment data for July and August, which recorded huge swings in the jobless rate. The ABS is also planning to review the methods it uses for seasonal adjustment to ensure its figures are as accurate as possible. Rob Hyndman, a professor of statistics at Monash University and member of the bureau's methodology advisory board, answers our questions:

A cautionary tale

```
{r abs1, echo=FALSE} employed <- tsibble(    Time =  
yearmonth("1978 Feb") + 0:439,    Employed = c(  
5985.7, 6040.6, 6054.2, 6038.3, 6031.3, 6036.1,  
6005.4, 6024.3, 6045.9, 6033.8, 6125.4, 5971.3,  
6050.7, 6096.2, 6087.7, 6075.6, 6095.7, 6103.9,  
6078.5, 6157.8, 6164.0, 6188.8, 6257.2, 6112.9,  
6207.2, 6278.7, 6224.9, 6273.4, 6269.9, 6314.1,  
6281.4, 6360.0, 6320.2, 6342.0, 6426.6, 6253.0,  
6356.5, 6428.1, 6426.3, 6412.4, 6413.9, 6425.3,  
6393.7, 6502.7, 6445.3, 6433.3, 6506.9, 6355.5,  
6432.4, 6497.4, 6431.6, 6440.9, 6414.3, 6425.9,
```

A cautionary tale

```
employed |>  
  autoplot(Employed) +  
  labs(title = "Total employed", y = "Thousands")
```

A cautionary tale

```
{r abs4, fig.height=2.8} employed |> filter(Year >= 2005) |> autoplot(Employed) + labs(title = "Total employed", y = "Thousands")
```

A cautionary tale

```
{r abs5, fig.height=2.8} employed |> filter(Year >= 2005) |> gg_season(Employed, labels = "right") + labs(title = "Total employed", y = "Thousands")
```

A cautionary tale

```
{r abs6, fig.height=2} employed |> mutate(diff =  
difference(Employed)) |> filter(Month == "Sep") |>  
ggplot(aes(y = diff, x = 1)) + geom_boxplot() +  
coord_flip() + labs(title = "Sep - Aug: total  
employed", y = "Thousands") +  
scale_x_continuous(breaks = NULL, labels = NULL)
```

A cautionary tale

```
{r abs7, fig.height=2.8} dcmp <- employed |>  
filter(Year >= 2005) |> model(stl = STL(Employed ~  
season(window = 11), robust = TRUE)) components(dcmp)  
|> autoplot()
```

A cautionary tale

```
{r abs8, fig.height=2.8} components(dcmp) |>  
filter(year(Time) == 2013) |> gg_season(season_year)  
+ labs(title = "Seasonal component") + guides(colour  
= "none")
```

A cautionary tale

```
components(dcmp) |>  
  as_tsibble() |>  
  autoplot(season_adjust)
```

A cautionary tale

- August 2014 employment numbers higher than expected.
- Supplementary survey usually conducted in August for employed people.
- Most likely, some employed people were claiming to be unemployed in August to avoid supplementary questions.
- Supplementary survey not run in 2014, so no motivation to lie about employment.
- In previous years, seasonal adjustment fixed the problem.
- The ABS has now adopted a new method to avoid the bias.

Outline

- 1 Time series data and tsibbles
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Lab Session 1
- 5 Time plots
- 6 Lab Session 2
- 7 Seasonal plots
- 8 Lab Session 3
- 9 Seasonal or cyclic?
- 10 Lag plots and autocorrelation
- 11 Lab Session 4
- 12 White noise
- 13 Lab Session 5
- 14 Per capita adjustments
- 15 Lab Session 6
- 16 Inflation adjustments
- 17 Mathematical transformations
- 18 Lab Session 7
- 19 Time series decompositions
- 20 Lab Session 8
- 21 Multiple seasonality
- 22 A cautionary tale
- 23 Notice: Material planned to change
- 24 STL Features
- 25 Lab Session 9
- 26 Dimension reduction for features

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL Features
- Lab Session 9
- Dimension reduction for features

Notice: Material planned to change

This material is planned to be updated to better align with the training needs of the Department of Education.

In particular, this section will be removed to make time for:

- econometric concepts
- multivariate models.

Outline



Time series data and tsibbles

Example: Australian prison population

Example: Australian pharmaceutical sales

Lab Session 1

Time plots

Lab Session 2

Seasonal plots

Lab Session 3

Seasonal or cyclic?

Lag plots and autocorrelation

Lab Session 4

White noise

Lab Session 5

Per capita adjustments

Lab Session 6

Inflation adjustments

Mathematical transformations

Lab Session 7

Time series decompositions

Lab Session 8

Multiple seasonality

A cautionary tale

Notice: Material planned to change

STL Features

Lab Session 9

Dimension reduction for features

Strength of seasonality and trend

STL decomposition

$$y_t = T_t + S_t + R_t$$

Seasonal strength

$$\max \left(0, 1 - \frac{\text{Var}(R_t)}{\text{Var}(S_t + R_t)} \right)$$

Trend strength

$$\max \left(0, 1 - \frac{\text{Var}(R_t)}{\text{Var}(T_t + R_t)} \right)$$

Feature extraction and statistics

```
tourism |> features(Trips, feat_stl)
```

Feature extraction and statistics

```
#| label: features-plot
#| fig-height: 3.6
tourism |>
  features(Trips, feat_stl) |>
  ggplot(aes(x = trend_strength, y = seasonal_strength_year, col = Purpose)) +
  geom_point() + facet_wrap(vars(State))
```

Feature extraction and statistics

```
#| label: features-plot
#| fig-height: 3.6
tourism |>
  features(Trips, feat_stl) |>
  ggplot(aes(x = trend_strength, y = seasonal_strength_year, col = Purpose)) +
  geom_point() + facet_wrap(vars(State))
```

- Holidays more seasonal than other travel.
- WA has strongest trends.

Feature extraction and statistics

Find the most seasonal time series:

```
most_seasonal <- tourism |>  
  features(Trips, feat_stl) |>  
  filter(seasonal_strength_year == max(seasonal_strength_year))
```

Feature extraction and statistics

Find the most seasonal time series:

```
most_seasonal <- tourism |>
  features(Trips, feat_stl) |>
  filter(seasonal_strength_year == max(seasonal_strength_year))

{r extreme2, fig.height=1.8} tourism |> right_join(most_seasonal, by = c("State",
"Region", "Purpose")) |> ggplot(aes(x = Quarter, y = Trips)) + geom_line() +
facet_grid(vars(State, Region, Purpose))
```

Feature extraction and statistics

Find the most trended time series:

```
most_trended <- tourism |>  
  features(Trips, feat_stl) |>  
  filter(trend_strength == max(trend_strength))
```

Feature extraction and statistics

Find the most trended time series:

```
most_trended <- tourism |>
  features(Trips, feat_stl) |>
  filter(trend_strength == max(trend_strength))
```

```
{r extreme4, fig.height=1.8} tourism |> right_join(most_trended, by = c("State",
"Region", "Purpose")) |> ggplot(aes(x = Quarter, y = Trips)) + geom_line() +
facet_grid(vars(State, Region, Purpose))
```

Outline



Time series data and tsibbles

Example: Australian prison population

Example: Australian pharmaceutical sales

Lab Session 1

Time plots

Lab Session 2

Seasonal plots

Lab Session 3

Seasonal or cyclic?

Lag plots and autocorrelation

Lab Session 4

White noise

Lab Session 5

Per capita adjustments

Lab Session 6

Inflation adjustments

Mathematical transformations

Lab Session 7

Time series decompositions

Lab Session 8

Multiple seasonality

A cautionary tale

Notice: Material planned to change

STL Features

Lab Session 9

Dimension reduction for features

Lab Session 9

- Use `GGally::ggpairs()` to look at the relationships between the STL-based features. You might wish to change `seasonal_peak_year` and `seasonal_trough_year` to factors.
- Which is the peak quarter for holidays in each state?

Feature extraction and statistics

```
tourism |> features(Trips, feat_acf)
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features
- Lab Session 10

Feature extraction and statistics

All features from the feasts package

```
#| echo: false
#| include: false
# Save pdf figures
savepdf <- function(file, width = 16, height = 10) {
  fname <- paste("figs/", file, ".pdf", sep = "")
  Cairo::CairoPDF(fname, width = width / 2.54, height = height / 2.54, pointsize = 12)
  par(mgp = c(2.2, 0.45, 0), tcl = -0.4, mar = c(3.3, 3.6, 1.1, 1.1))
}
endpdf <- function() {
  crop::dev.off.crop(fname)
}
# Compute features
tourism_features <- tourism |>
  features(Trips, feature_set(pkgs = "feasts"))
# Compute PCs
pcs <- tourism_features |>
  select(-State, -Region, -Purpose) |>
  prcomp(scale = TRUE) |>
  broom::augment(tourism_features)
```

Feature extraction and statistics

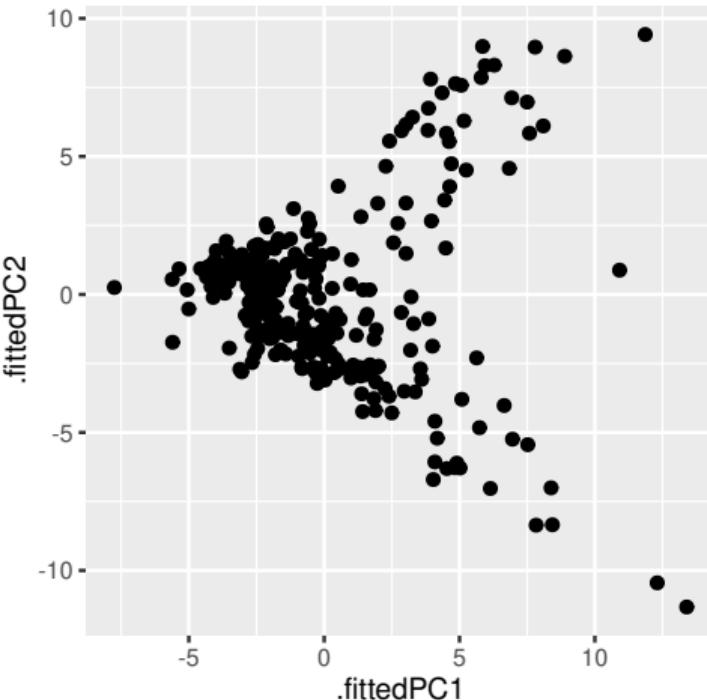
```
pcs <- tourism_features |>  
  select(-State, -Region, -Purpose) |>  
  prcomp(scale = TRUE) |>  
  broom::augment(tourism_features)  
{r echo=FALSE} pcs
```

Principal components based
on all features from the
feasts package

Feature extraction and statistics

```
pcs |> ggplot(aes(x=.fittedPC1, y=.fittedPC2)) +  
  geom_point() + theme(aspect.ratio=1)
```

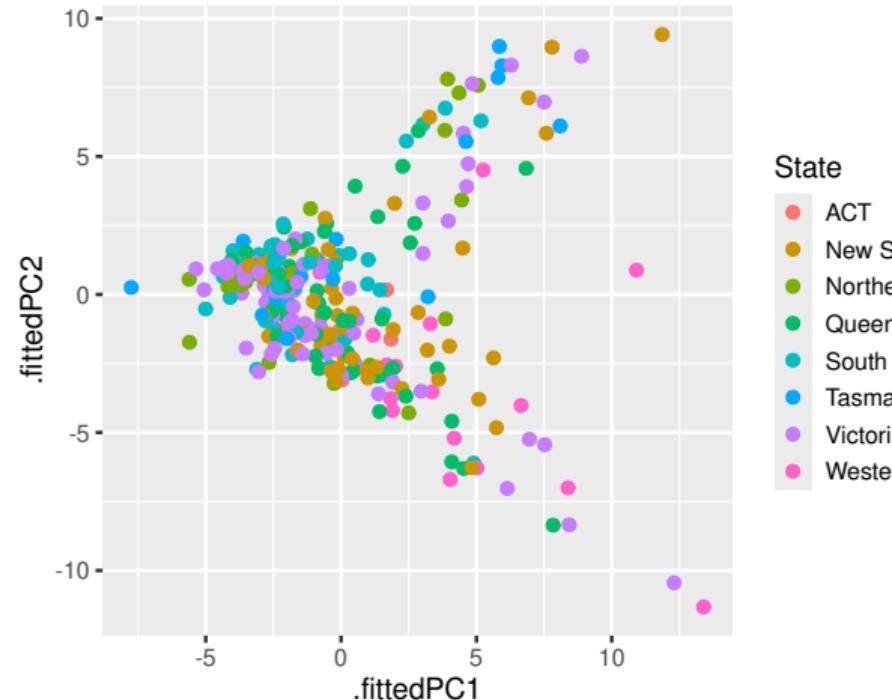
Principal components based on all features from the feasts package



Feature extraction and statistics

```
pcs |> ggplot(aes(x=.fittedPC1, y=.fittedPC2, col=State)) +  
  geom_point() + theme(aspect.ratio=1)
```

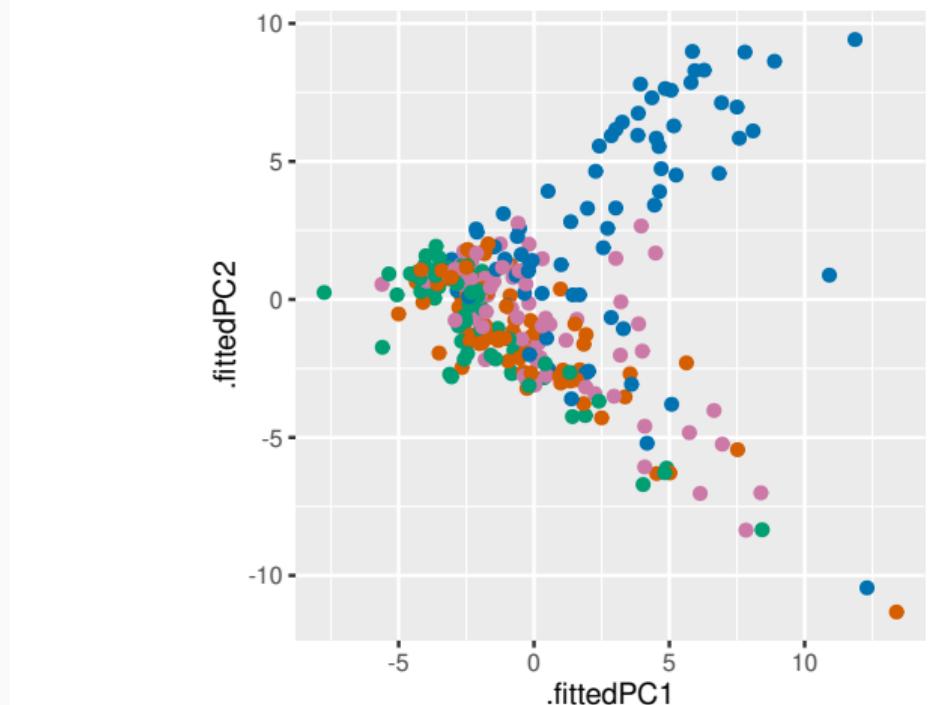
Principal components based on all features from the feasts package



Feature extraction and statistics

```
pcs |> ggplot(aes(x=.fittedPC1, y=.fittedPC2, col=Purpose)) +  
  geom_point() + theme(aspect.ratio=1)
```

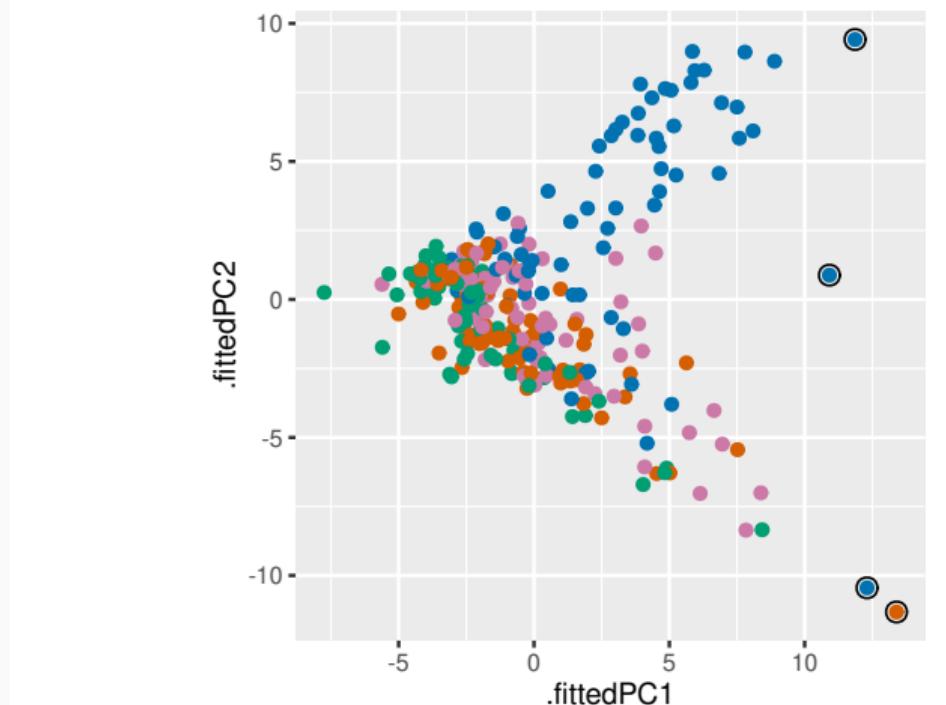
Principal components based on all features from the feasts package



Feature extraction and statistics

```
pcs |> ggplot(aes(x=.fittedPC1, y=.fittedPC2, col=Purpose)) +  
  geom_point() + theme(aspect.ratio=1)
```

Principal components based on all features from the feasts package



Feature extraction and statistics

```
{r outliers2, fig.height=4.5, fig.width=12, out.height="55%"} outliers |> left_join(tourism, by = c("State", "Region", "Purpose")) |> mutate(Series = glue("{State}", "{Region}", "{Purpose}", .sep = "\n\n")) |> ggplot(aes(x = Quarter, y = Trips)) + geom_line() + facet_grid(Series ~ .) + labs(title = "Outlying time series in PC space")
```

Outline



Time series data and tsibbles

Example: Australian prison population

Example: Australian pharmaceutical sales

Lab Session 1

Time plots

Lab Session 2

Seasonal plots

Lab Session 3

Seasonal or cyclic?

Lag plots and autocorrelation

Lab Session 4

White noise

Lab Session 5

Per capita adjustments

Lab Session 6

Inflation adjustments

Mathematical transformations

Lab Session 7

Time series decompositions

Lab Session 8

Multiple seasonality

A cautionary tale

Notice: Material planned to change

STL features

Lab Session 9

Dimension reduction for features

Lab Session 10

Lab Session 10

- Use a feature-based approach to look for outlying series in PBS.
- What is unusual about the series you identify as outliers?

Outline

- 1 Time series data and tsibbles
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Lab Session 1
- 5 Time plots
- 6 Lab Session 2
- 7 Seasonal plots
- 8 Lab Session 3
- 9 Seasonal or cyclic?
- 10 Lag plots and autocorrelation
- 11 Lab Session 4
- 12 White noise
- 13 Lab Session 5
- 14 Per capita adjustments
- 15 Lab Session 6
- 16 Inflation adjustments
- 17 Mathematical transformations
- 18 Lab Session 7
- 19 Time series decompositions
- 20 Lab Session 8
- 21 Multiple seasonality
- 22 A cautionary tale
- 23 Notice: Material planned to change
- 24 STL Features
- 25 Lab Session 9
- 26 Dimension reduction for features

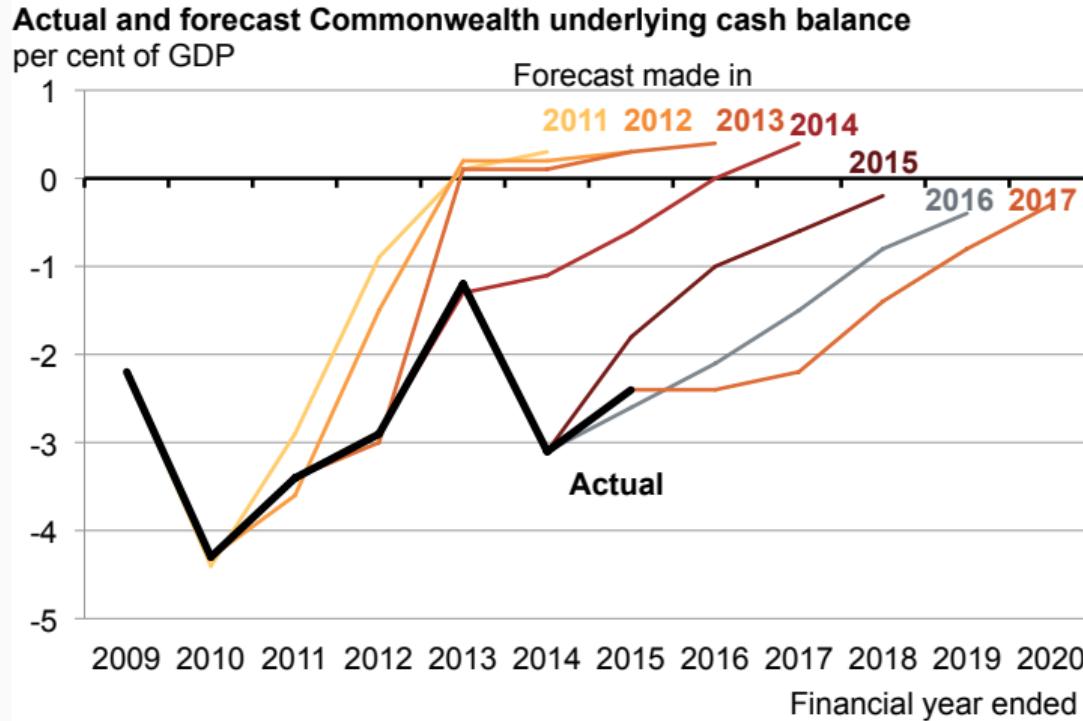
Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Forecasting is difficult

Commonwealth plans to drift back to surplus
show the triumph of experience over hope

GRATTAN
Institute



What can we forecast?



What can we forecast?



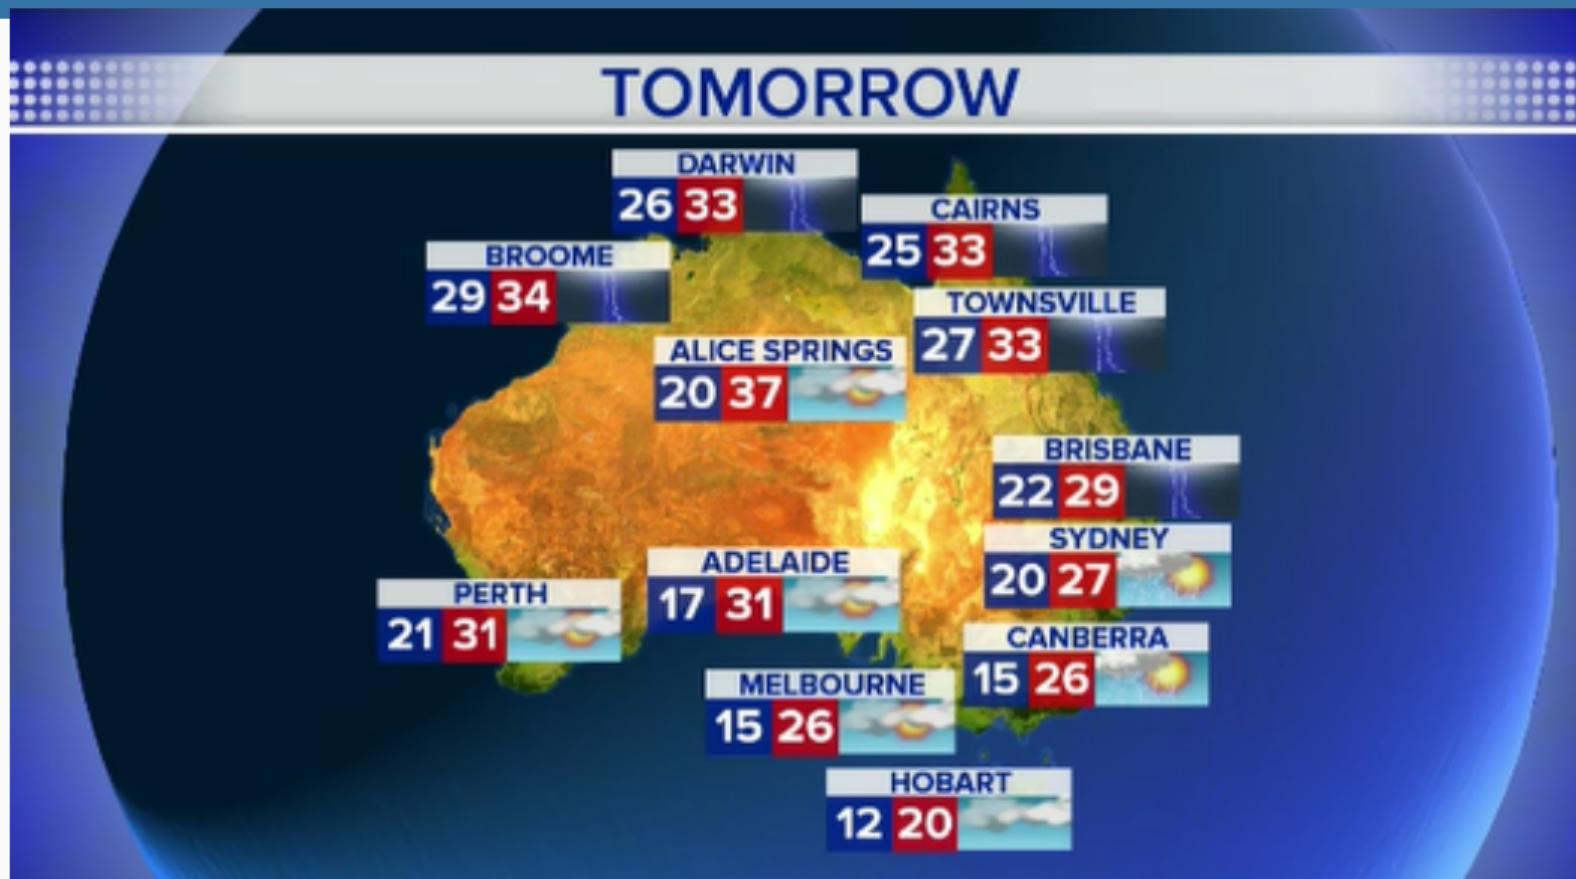
What can we forecast?



What can we forecast?



What can we forecast?



What can we forecast?



What can we forecast?



Which is easiest to forecast?

- 1 daily electricity demand in 3 days time
- 2 timing of next Halley's comet appearance
- 3 time of sunrise this day next year
- 4 Google stock price tomorrow
- 5 Google stock price in 6 months time
- 6 maximum temperature tomorrow
- 7 exchange rate of \$US/AUS next week
- 8 total sales of drugs in Australian pharmacies next month

Which is easiest to forecast?

- 1 daily electricity demand in 3 days time
 - 2 timing of next Halley's comet appearance
 - 3 time of sunrise this day next year
 - 4 Google stock price tomorrow
 - 5 Google stock price in 6 months time
 - 6 maximum temperature tomorrow
 - 7 exchange rate of \$US/AUS next week
 - 8 total sales of drugs in Australian pharmacies next month
-
- how do we measure “easiest”?
 - what makes something easy/difficult to forecast?

Factors affecting forecastability

Something is easier to forecast if:

- we have a good understanding of the factors that contribute to it
- there is lots of data available;
- the forecasts cannot affect the thing we are trying to forecast.
- there is relatively low natural/unexplainable random variation.
- the future is somewhat similar to the past

Random futures

```
{r austas, echo=FALSE} # Grab ABS data austas <-
readxl::read_excel("data/340101.xlsx", sheet =
"Data1", skip = 9) |>    rename(date = `Series ID` ,
value = A85375847A) |>    select(date, value) |>
transmute(      Month = yearmonth(date),      Visitors =
value / 1e3      ) |>    bind_rows(tibble(      Month =
yearmonth(seq(as.Date("2021-11-01"), by = "1 month",
length = 2)),      Visitors = NA_real_      )) |>
as_tsibble(index = Month) |>    filter(Month >=
yearmonth("2000 Jan")) # Fit ETS model fit <- austas |>
filter(Month < yearmonth("2018 Jan")) |>
```

Random futures

```
{r austas, echo=FALSE} # Grab ABS data austas <-
readxl::read_excel("data/340101.xlsx", sheet =
"Data1", skip = 9) |>    rename(date = `Series ID` ,
value = A85375847A) |>    select(date, value) |>
transmute(      Month = yearmonth(date),      Visitors =
value / 1e3      ) |>    bind_rows(tibble(      Month =
yearmonth(seq(as.Date("2021-11-01"), by = "1 month",
length = 2)),      Visitors = NA_real_      )) |>
as_tsibble(index = Month) |>    filter(Month >=
yearmonth("2000 Jan")) # Fit ETS model fit <- austas |>
filter(Month < yearmonth("2018 Jan")) |>
```

Random futures

A forecast is an estimate of the probabilities of possible futures.

```
{r austaa2, dependson='austaa', echo=FALSE}  
aligned_plots[[2]]
```

Simulated futures
from an ETS model

Random futures

A forecast is an estimate of the probabilities of possible futures.

```
{r austaa3, dependson='austaa', echo=FALSE}  
aligned_plots[[3]]
```

Simulated futures
from an ETS model

Random futures

A forecast is an estimate of the probabilities of possible futures.

```
{r aust4, dependson='aust', echo=FALSE}  
aligned_plots[[4]]
```

Simulated futures
from an ETS model

Random futures

A forecast is an estimate of the probabilities of possible futures.

```
{r austas5, dependson='austas', echo=FALSE}  
aligned_plots[[5]]
```

Simulated futures
from an ETS model

Random futures

A forecast is an estimate of the probabilities of possible futures.

```
{r austaa6, dependson='austaa', echo=FALSE}  
aligned_plots[[6]]
```

Simulated futures
from an ETS model

Random futures

A forecast is an estimate of the probabilities of possible futures.

```
{r austat7, dependson='austat', echo=FALSE}  
aligned_plots[[7]]
```

Simulated futures
from an ETS model

Random futures

A forecast is an estimate of the probabilities of possible futures.

```
{r austaa, dependson='austa', echo=FALSE}  
aligned_plots[[8]]
```

Simulated futures
from an ETS model

Random futures

A forecast is an estimate of the probabilities of possible futures.

```
{r aust9, dependson='austa', echo=FALSE}  
aligned_plots[[9]]
```

Simulated futures
from an ETS model

Random futures

A forecast is an estimate of the probabilities of possible futures.

```
{r aust9b, dependson='austa', echo=FALSE}
aligned_plots[[9]] + geom_line( colour =
"black", data = aust9b |> filter(Month >=
yearmonth("2018 Jan")) |> mutate(Month =
as.Date(Month)) )
```

Simulated futures
from an ETS model

Random futures

A forecast is an estimate of the probabilities of possible futures.

```
{r austab9, dependson='austa', echo=FALSE}  
aligned_plots[[9]] + geom_line( colour =  
"black", data = austab |> filter(Month >=  
yearmonth("2018 Jan")) |> mutate(Month =  
as.Date(Month)) )
```

“He who sees the past as surprise-free is
bound to have a future full of surprises.”

(Amos Tversky)

Simulated futures
from an ETS model

Statistical forecasting

- Thing to be forecast: y_{T+h} .
- What we know: y_1, \dots, y_T .
- Forecast distribution: $y_{T+h|t} = y_{T+h} \mid \{y_1, y_2, \dots, y_T\}$.
- Point forecast: $\hat{y}_{T+h|T} = E[y_{T+h} \mid y_1, \dots, y_T]$.
- Forecast variance: $\text{Var}[y_t \mid y_1, \dots, y_T]$
- Prediction interval is a range of values of y_{T+h} with high probability.

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Some simple forecasting methods

```
{r ausbeer, fig.height=4.6, echo=FALSE} new_production  
  <- aus_production |> filter(year(Quarter) >= 1992)  
new_production |> autoplot(Beer) + labs( x =  
  "Year", y = "megalitres", title = "Australian  
quarterly beer production" )
```

How would you forecast these series?

Some simple forecasting methods

```
{r pigs, fig.height=4.6, echo=FALSE} aus_livestock |>  
filter(    between(year(Month), 1992, 1996),  
Animal == "Pigs", State == "Victoria" ) |>  
autoplot(Count) +  labs(    x = "Year", y =  
"thousands",    title = "Number of pigs slaughtered  
in Victoria, 1990-1995" )
```

How would you forecast these series?

Some simple forecasting methods

```
{r dj, fig.height=4.6, echo=FALSE} gafa_stock |>  
filter(Symbol == "FB", Date >= ymd("2018-01-01")) |>  
autoplot(Close) + labs( title = "Facebook  
closing stock price in 2018", x = "Date", y =  
"Closing price ($USD)" )
```

How would you forecast these series?

Some simple forecasting methods

MEAN(y): Average method

- Forecast of all future values is equal to mean of historical data $\{y_1, \dots, y_T\}$.
- Forecasts: $\hat{y}_{T+h|T} = \bar{y} = (y_1 + \dots + y_T)/T$

```
“{r mean-method-explained, echo=FALSE, message=FALSE,  
warning=FALSE, fig.height = 3.4} bricks <- aus_production |>  
filter(!is.na(Bricks)) |> mutate(average = mean(Bricks))  
fc <- bricks |> model(MEAN(Bricks)) |> forecast(h = “5 years”)  
bricks |> ggplot(aes(x = Quarter, y = Bricks)) + geom_line() +  
geom_line(aes(y = average), colour = “blue”, linetype = “dashed”) +  
geom_line(aes(y = .mean), data = fc, colour = “blue”) + labs(title =  
“Clay brick production in Australia”)
```

Some simple forecasting methods

SNAIVE($y \sim \text{lag}(m)$): Seasonal naïve method

- Forecasts equal to last value from same season.
- Forecasts: $\hat{y}_{T+h|T} = y_{T+h-m(k+1)}$, where m = seasonal period and k is the integer part of $(h - 1)/m$.

```
{r snaive-method-explained, echo = FALSE, warning = FALSE,
fig.height = 3.4} bricks |> model(SNAIVE(Bricks ~
lag("year"))) |> forecast(h = "5 years") |>
autoplot(filter(bricks, year(Quarter) > 1990), level = NULL)
+ geom_point(aes(y = Bricks), data = slice(bricks, (n() -
3):n()), colour = "blue") + labs(title = "Clay brick
production in Australia")
```

Some simple forecasting methods

RW(y ~ drift()): Drift method

- Forecasts equal to last value plus average change.
- Forecasts:

$$\begin{aligned}\hat{y}_{T+h|T} &= y_T + \frac{h}{T-1} \sum_{t=2}^T (y_t - y_{t-1}) \\ &= y_T + \frac{h}{T-1} (y_T - y_1).\end{aligned}$$

- Equivalent to extrapolating a line drawn between first and last observations.

Some simple forecasting methods

Drift method

```
{r drift-method-explained, echo = FALSE, warning =  
FALSE} aus_production |> filter(!is.na(Bricks)) |>  
model(RW(Bricks ~ drift())) |> forecast(h = "5  
years") |> autoplot(aus_production, level = NULL) +  
geom_line(aes(y = Bricks), data =  
slice(aus_production, range(cumsum(!is.na(Bricks)))),  
linetype = "dashed", colour = "blue") +  
labs(title = "Clay brick production in Australia")
```

Model fitting

The `model()` function trains models to data.

```
brick_fit <- aus_production |>  
  filter(!is.na(Bricks)) |>  
  model(  
    `Seasonal_naïve` = SNAIVE(Bricks),  
    `Naïve` = NAIVE(Bricks),  
    Drift = RW(Bricks ~ drift()),  
    Mean = MEAN(Bricks)  
  )
```

```
{r brick-model2, echo=FALSE, dependson='brick-model'} brick_fit
```

A `mable` is a model table, each cell corresponds to a fitted model.

Producing forecasts

```
{r brick-fc, echo = TRUE, dependson='brick-model'} brick_fc <- brick_fit |>  
forecast(h = "5 years")  
  
{r brick-fbl, echo = FALSE, dependson='brick-fc'} print(brick_fc, n = 4)
```

A fable is a forecast table with point forecasts and distributions.

Visualising forecasts

```
{r brick-fc-plot, warning=FALSE, message=FALSE, fig.height=3.4,
dependson='brick-fc'} brick_fc |> autoplot(aus_production, level = NULL) +
labs(title = "Forecasts for quarterly clay brick production", x =
"Year", y = "Millions of bricks") + guides(colour = guide_legend(title =
"Forecast"))
```

Prediction intervals

```
{r brick-fc-interval, dependson='brick-fc'} brick_fc |> hilo(level = c(50, 75))
```

Prediction intervals

```
{r brick-fc-interval2, dependson='brick-fc'} brick_fc |> hilo(level = c(50, 75)) |> mutate(lower = `50%`$lower, upper = `50%`$upper)
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL Features
- Lab Session 9
- Dimension reduction for features

Lab Session 11

- Produce forecasts using an appropriate benchmark method for student enrolments in Australia (as shown yesterday). Plot the results using `autoplot()`.
- Produce forecasts using an appropriate benchmark method for total Australian retail turnover (aggregate `aus_retail`). Plot the results using `autoplot()`.

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Fitted values

- $\hat{y}_{t|t-1}$ is the forecast of y_t based on observations y_1, \dots, y_{t-1} .
- We call these “fitted values”.
- Sometimes drop the subscript: $\hat{y}_t \equiv \hat{y}_{t|t-1}$.
- Often not true forecasts since parameters are estimated on all data.

For example:

- $\hat{y}_t = \bar{y}$ for average method.
- $\hat{y}_t = y_{t-1} + (y_T - y_1)/(T - 1)$ for drift method.

Forecasting residuals

Residuals in forecasting: difference between observed value and its fitted value: $e_t = y_t - \hat{y}_{t|t-1}$.

Forecasting residuals

Residuals in forecasting: difference between observed value and its fitted value: $e_t = y_t - \hat{y}_{t|t-1}$.

Assumptions

- 1 $\{e_t\}$ uncorrelated. If they aren't, then information left in residuals that should be used in computing forecasts.
- 2 $\{e_t\}$ have mean zero. If they don't, then forecasts are biased.

Forecasting residuals

Residuals in forecasting: difference between observed value and its fitted value: $e_t = y_t - \hat{y}_{t|t-1}$.

Assumptions

- 1 $\{e_t\}$ uncorrelated. If they aren't, then information left in residuals that should be used in computing forecasts.
- 2 $\{e_t\}$ have mean zero. If they don't, then forecasts are biased.

Useful properties (for prediction intervals)

- 3 $\{e_t\}$ have constant variance.
- 4 $\{e_t\}$ are normally distributed.

Facebook closing stock price

```
fb_stock <- gafa_stock |>  
  filter(Symbol == "FB")  
fb_stock |> autoplot(Close)
```

Facebook closing stock price

```
fb_stock <- fb_stock |>  
  mutate(trading_day = row_number()) |>  
  update_tsibble(index = trading_day, regular = TRUE)  
fit <- fb_stock |> model(NAIVE(Close))  
augment(fit)
```

Facebook closing stock price

```
{r dj4, echo=TRUE, warning=FALSE, fig.height=3.4, dependson="augment"}  
augment(fit) |> ggplot(aes(x = trading_day)) + geom_line(aes(y = Close,  
colour = "Data")) + geom_line(aes(y = .fitted, colour = "Fitted"))
```

Facebook closing stock price

```
{r dj4a, echo=TRUE, warning=FALSE, fig.height=3.4, dependson="augment"}  
augment(fit) |> filter(trading_day > 1100) |> ggplot(aes(x = trading_day))  
+ geom_line(aes(y = Close, colour = "Data")) + geom_line(aes(y = .fitted,  
colour = "Fitted"))
```

Facebook closing stock price

```
{r dj5, echo=TRUE, warning = FALSE, dependson="augment"} augment(fit) |>  
autoplot(.resid) + labs(x = "Day", y = "", title = "Residuals from naïve  
method")
```

Facebook closing stock price

```
{r dj6, warning=FALSE, fig.height=3.4, dependson="augment"}  
augment(fit) |> ggplot(aes(x = .resid)) + geom_histogram(bins =  
150) + labs(title = "Histogram of residuals")
```

Facebook closing stock price

```
{r dj7, dependson="augment"} augment(fit) |> ACF(.resid) |>  
autoplot() + labs(title = "ACF of residuals")
```

ACF of residuals

- We assume that the residuals are white noise (uncorrelated, mean zero, constant variance). If they aren't, then there is information left in the residuals that should be used in computing forecasts.
- So a standard residual diagnostic is to check the ACF of the residuals of a forecasting method.
- We *expect* these to look like white noise.

Combined diagnostic graph

```
{r dj8, dependson="augment"} fit |> gg_tsresiduals()
```

Ljung-Box test

Test whether *whole set* of r_k values is significantly different from zero set.

$$Q = T(T + 2) \sum_{k=1}^{\ell} (T - k)^{-1} r_k^2 \quad \text{where } \ell = \text{max lag and } T = \# \text{ observations}$$

- If each r_k close to zero, Q will be **small**.
- If some r_k values large (+ or -), Q will be **large**.
- My preferences: $h = 10$ for non-seasonal data, $h = 2m$ for seasonal data.
- If data are WN and T large, $Q \sim \chi^2$ with ℓ degrees of freedom.

Ljung-Box test

$$Q = T(T + 2) \sum_{k=1}^{\ell} (T - k)^{-1} r_k^2 \quad \text{where } \ell = \text{max lag and } T = \# \text{ observations.}$$

```
{r dj9extra, echo=FALSE, fig.height=1.65} augment(fit) |> ACF(.resid,  
lag_max = 10) |> autoplot() + labs(title = "ACF of residuals")  
{r dj9, echo=TRUE, dependson="augment"} # lag = h augment(fit) |>  
features(.resid, ljung_box, lag = 10)
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Lab Session 12

- Compute RW w/ drift forecasts for total student enrolments in Australia.
- Test if the residuals are white noise. What do you conclude?

Outline



Time series data and tsibbles

Example: Australian prison population

Example: Australian pharmaceutical sales

Lab Session 1

Time plots

Lab Session 2

Seasonal plots

Lab Session 3

Seasonal or cyclic?

Lag plots and autocorrelation

Lab Session 4

White noise

Lab Session 5

Per capita adjustments

Lab Session 6

Inflation adjustments

Mathematical transformations

Lab Session 7

Time series decompositions

Lab Session 8

Multiple seasonality

A cautionary tale

Notice: Material planned to change

STL features

Lab Session 9

Dimension reduction for features

Training and test sets

```
{r traintest, fig.height=1, echo=FALSE} train <- 1:18  
test <- 19:24 par(mar = c(0, 0, 0, 0)) plot(0, 0, xlim  
= c(0, 26), ylim = c(0, 2), xaxt = "n", yaxt = "n",  
bty = "n", xlab = "", ylab = "", type = "n") arrows(0,  
0.5, 25, 0.5, 0.05) points(train, train * 0 + 0.5, pch  
= 19, col = "blue") points(test, test * 0 + 0.5, pch =  
19, col = "red") text(26, 0.5, "time") text(10, 1,  
"Training data", col = "blue") text(21, 1, "Test  
data", col = "red")
```

- A model which fits the training data well will not necessarily forecast well.

Measures of forecast accuracy

```
beer_fit <- aus_production |>
  filter(between(year(Quarter), 1992, 2007)) |>
  model(
    snaive = SNAIVE(Beer),
    mean = MEAN(Beer)
  )
beer_fit |>
  forecast(h = "3 years") |>
  autoplot(aus_production, level = NULL) +
  labs(title ="Forecasts for quarterly beer production",
       x ="Year", y ="Megalitres") +
  guides(colour = guide_legend(title = "Forecast"))
```

Measures of forecast accuracy

```
{r beer-fc-1, echo=FALSE, fig.height=4} beer_fit <-  
aus_production |> filter(between(year(Quarter),  
1992, 2007)) |> model(    snaive = SNAIVE(Beer),  
mean = MEAN(Beer)    ) beer_fit |> forecast(h = "3  
years") |> autoplot(aus_production, level = NULL) +  
labs(    title = "Forecasts for quarterly beer  
production",    x = "Year", y = "Megalitres"    ) +  
guides(colour = guide_legend(title = "Forecast"))
```

Measures of forecast accuracy

y_{T+h} = $(T + h)$ th observation, $h = 1, \dots, H$

$\hat{y}_{T+h|T}$ = its forecast based on data up to time T .

$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T}$

MAE = $\text{mean}(|e_{T+h}|)$

MSE = $\text{mean}(e_{T+h}^2)$

RMSE = $\sqrt{\text{mean}(e_{T+h}^2)}$

MAPE = $100\text{mean}(|e_{T+h}|/|y_{T+h}|)$

Measures of forecast accuracy

y_{T+h} = $(T + h)$ th observation, $h = 1, \dots, H$

$\hat{y}_{T+h|T}$ = its forecast based on data up to time T .

$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T}$

MAE = $\text{mean}(|e_{T+h}|)$

MSE = $\text{mean}(e_{T+h}^2)$

RMSE = $\sqrt{\text{mean}(e_{T+h}^2)}$

MAPE = $100\text{mean}(|e_{T+h}|/|y_{T+h}|)$

- MAE, MSE, RMSE are all scale dependent.
- MAPE is scale independent but is only sensible if $y_t \gg 0$ for all t , and y has a natural zero.

Measures of forecast accuracy

Mean Absolute Scaled Error

$$\text{MASE} = \text{mean}(|e_{T+h}|/Q)$$

Measures of forecast accuracy

Mean Absolute Scaled Error

$$\text{MASE} = \text{mean}(|e_{T+h}| / Q)$$

- For non-seasonal series, scale uses naïve forecasts:

$$Q = \frac{1}{T-1} \sum_{t=2}^T |y_t - y_{t-1}|$$

- For seasonal series, scale uses seasonal naïve forecasts:

$$Q = \frac{1}{T-m} \sum_{t=m+1}^T |y_t - y_{t-m}|$$

where m is the seasonal frequency

Measures of forecast accuracy

Mean Absolute Scaled Error

$$\text{MASE} = \text{mean}(|e_{T+h}| / Q)$$

- For non-seasonal series, scale uses naïve forecasts:

$$Q = \frac{1}{T-1} \sum_{t=2}^T |y_t - y_{t-1}|$$

- For seasonal series, scale uses seasonal naïve forecasts:

$$Q = \frac{1}{T-m} \sum_{t=m+1}^T |y_t - y_{t-m}|$$

where m is the seasonal frequency

Proposed by Hyndman and Koehler (IJF, 2006).

Measures of forecast accuracy

Root Mean Squared Scaled Error

$$\text{RMSSE} = \sqrt{\text{mean}(e_{T+h}^2/Q)}$$

- For non-seasonal series, scale uses naïve forecasts:

$$Q = \frac{1}{T-1} \sum_{t=2}^T (y_t - y_{t-1})^2$$

- For seasonal series, scale uses seasonal naïve forecasts:

$$Q = \frac{1}{T-m} \sum_{t=m+1}^T (y_t - y_{t-m})^2$$

where m is the seasonal frequency

Proposed by Hyndman and Koehler (IJF, 2006).

Measures of forecast accuracy

```
{r beer-test-accuracy, dependson='beer-fc-1'} beer_fc <- forecast(beer_fit, h =  
"3 years") accuracy(beer_fc, aus_production)
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Lab Session 13

- Create a training set for Australian student enrolments by withholding the last four years as a test set.
- Fit all the appropriate benchmark methods to the training set and forecast the periods covered by the test set.
- Compute the accuracy of your forecasts. Which method does best?
- Repeat the exercise using the Australian takeaway food turnover data (`aus_retail`) with a test set of four years.

Outline

- 1 Time series data and tsibbles
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Lab Session 1
- 5 Time plots
- 6 Lab Session 2
- 7 Seasonal plots
- 8 Lab Session 3
- 9 Seasonal or cyclic?
- 10 Lag plots and autocorrelation
- 11 Lab Session 4
- 12 White noise
- 13 Lab Session 5
- 14 Per capita adjustments
- 15 Lab Session 6
- 16 Inflation adjustments
- 17 Mathematical transformations
- 18 Lab Session 7
- 19 Time series decompositions
- 20 Lab Session 8
- 21 Multiple seasonality
- 22 A cautionary tale
- 23 Notice: Material planned to change
- 24 STL features
- 25 Lab Session 9
- 26 Dimension reduction for features

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Pharmaceutical Benefits Scheme



Pharmaceutical Benefits Scheme

The Pharmaceutical Benefits Scheme (PBS) is the Australian government drugs subsidy scheme.

- Many drugs bought from pharmacies are subsidised to allow more equitable access to modern drugs.
- The cost to government is determined by the number and types of drugs purchased. Currently nearly 1% of GDP.
- The total cost is budgeted based on forecasts of drug usage.

Pharmaceutical Benefits Scheme



ABC News Online
AUSTRALIAN BROADCASTING CORPORATION

Select a Topic
from the list below

- ▶ [Top Stories](#)
- ▶ [Just In](#)
- ▶ [World](#)
- ▶ [Asia-Pacific](#)
- ▶ [Business](#)
- ▶ [Sport](#)
- ▶ [Arts](#)
- ▶ [Sci Tech](#)
- ▶ [Indigenous](#)
- ▶ [Weather](#)
- ▶ [Rural](#)
- ▶ [Local News](#)
- ▶ [Broadband](#)

SPECIALS

- ▶ [Federal Election](#)

Click "Refresh" or "Reload"
on your browser for the latest edition.

POLITICS

Opp demands drug price restriction after PBS budget blow-out

The Federal Opposition has called for tighter controls on drug prices after the Pharmaceutical Benefits Scheme (PBS) budget blew out by almost \$800 million.

The money was spent on two new drugs including the controversial anti-smoking aid Zyban, which dropped in price from \$220 to \$22 after it was listed on the PBS.



NewsRadio
Streaming audio news
LISTEN: WMP | Real

This Bulletin: Wed, May 30 2001 6:22 PM AEST

the Public Record

For full election coverage

FEATURES

Public Record

Federal Election 2001

[For a fresh perspective on the federal election, reach into ABC Online's campaign weblog, The Poll Vault.](#)

[Audio News Online](#)

Pharmaceutical Benefits Scheme

- In 2001: \$4.5 billion budget, under-forecasted by \$800 million.
- Thousands of products. Seasonal demand.
- Subject to covert marketing, volatile products, uncontrollable expenditure.
- Although monthly data available for 10 years, data are aggregated to annual values, and only the first three years are used in estimating the forecasts.
- All forecasts being done with the FORECAST function in MS-Excel!

Historical perspective

- Developed in the 1950s and 1960s as methods (algorithms) to produce point forecasts.
- Combine a “level”, “trend” (slope) and “seasonal” component to describe a time series.
- The rate of change of the components are controlled by “smoothing parameters”: α , β and γ respectively.
- Need to choose best values for the smoothing parameters (and initial states).
- Equivalent ETS state space models developed in the 1990s and 2000s.

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

Additively?

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

Additively?

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

Multiplicatively?

$$y_t = \ell_{t-1} b_{t-1} s_{t-m} (1 + \varepsilon_t)$$

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

Additively?

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

Multiplicatively?

$$y_t = \ell_{t-1} b_{t-1} s_{t-m} (1 + \varepsilon_t)$$

Perhaps a mix of both?

$$y_t = (\ell_{t-1} + b_{t-1}) s_{t-m} + \varepsilon_t$$

A model for levels, trends, and seasonalities

We want a model that captures the level (ℓ_t), trend (b_t) and seasonality (s_t).

How do we combine these elements?

Additively?

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

Multiplicatively?

$$y_t = \ell_{t-1} b_{t-1} s_{t-m} (1 + \varepsilon_t)$$

Perhaps a mix of both?

$$y_t = (\ell_{t-1} + b_{t-1}) s_{t-m} + \varepsilon_t$$

How do the level, trend and seasonal components evolve over time?

ETS models

General notation $\overset{\nearrow}{E} \overset{\uparrow}{T} \overset{\nwarrow}{S} : \text{ExponenTial Smoothing}$
Error Trend Season

Error: Additive ("A") or multiplicative ("M")

ETS models

General notation $\overset{\nearrow}{E} \overset{\uparrow}{T} \overset{\nwarrow}{S} : \text{ExponenTial Smoothing}$
Error Trend Season

Error: Additive ("A") or multiplicative ("M")

Trend: None ("N"), additive ("A"), multiplicative ("M"), or damped ("Ad" or "Md").

ETS models

General notation  **E T S : ExponenTial Smoothing**
Error Trend Season

Error: Additive ("A") or multiplicative ("M")

Trend: None ("N"), additive ("A"), multiplicative ("M"), or damped ("Ad" or "Md").

Seasonality: None ("N"), additive ("A") or multiplicative ("M")

ETS(A,N,N): SES with additive errors

Forecast equation

$$\hat{y}_{T+h|T} = \ell_T$$

Measurement equation

$$y_t = \ell_{t-1} + \varepsilon_t$$

State equation

$$\ell_t = \ell_{t-1} + \alpha \varepsilon_t$$

where $\varepsilon_t \sim \text{NID}(0, \sigma^2)$.

ETS(A,N,N): SES with additive errors

Forecast equation

$$\hat{y}_{T+h|T} = \ell_T$$

Measurement equation

$$y_t = \ell_{t-1} + \varepsilon_t$$

State equation

$$\ell_t = \ell_{t-1} + \alpha \varepsilon_t$$

where $\varepsilon_t \sim \text{NID}(0, \sigma^2)$.

- “innovations” or “single source of error” because equations have the same error process, ε_t .
- Measurement equation: relationship between observations and states.
- Transition/state equation(s): evolution of state(s) over time.

ETS(M,N,N): SES with multiplicative errors

Forecast equation

$$\hat{y}_{T+h|T} = \ell_T$$

Measurement equation

$$y_t = \ell_{t-1}(1 + \varepsilon_t)$$

State equation

$$\ell_t = \ell_{t-1}(1 + \alpha \varepsilon_t)$$

where $\varepsilon_t \sim \text{NID}(0, \sigma^2)$.

ETS(M,N,N): SES with multiplicative errors

Forecast equation

$$\hat{y}_{T+h|T} = \ell_T$$

Measurement equation

$$y_t = \ell_{t-1}(1 + \varepsilon_t)$$

State equation

$$\ell_t = \ell_{t-1}(1 + \alpha \varepsilon_t)$$

where $\varepsilon_t \sim \text{NID}(0, \sigma^2)$.

- Models with additive and multiplicative errors with the same parameters generate the same point forecasts but different prediction intervals.

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Holt's linear trend

Additive errors: ETS(A,A,N)

Forecast equation

$$\hat{y}_{T+h|T} = \ell_T + h b_T$$

Measurement equation

$$y_t = \ell_{t-1} + b_{t-1} + \varepsilon_t$$

State equations

$$\ell_t = \ell_{t-1} + b_{t-1} + \alpha \varepsilon_t$$

$$b_t = b_{t-1} + \beta \varepsilon_t$$

Holt's linear trend

Additive errors: ETS(A,A,N)

Forecast equation

$$\hat{y}_{T+h|T} = \ell_T + h b_T$$

Measurement equation

$$y_t = \ell_{t-1} + b_{t-1} + \varepsilon_t$$

State equations

$$\ell_t = \ell_{t-1} + b_{t-1} + \alpha \varepsilon_t$$

$$b_t = b_{t-1} + \beta \varepsilon_t$$

Multiplicative errors: ETS(M,A,N)

Forecast equation

$$\hat{y}_{T+h|T} = \ell_T + h b_T$$

Measurement equation

$$y_t = (\ell_{t-1} + b_{t-1})(1 + \varepsilon_t)$$

State equations

$$\ell_t = (\ell_{t-1} + b_{t-1})(1 + \alpha \varepsilon_t)$$

$$b_t = b_{t-1} + \beta(\ell_{t-1} + b_{t-1})\varepsilon_t$$

Example: Australian population

```
{r holt-fit, echo=TRUE} aus_economy <- global_economy |> filter(Country ==  
"Australia") |> mutate(Pop = Population / 1e6) fit <- aus_economy |> model(AAN =  
ETS(Pop)) report(fit)
```

Example: Australian population

```
{r holt-cmp, echo=TRUE, dependson='holt-fit'} components(fit)
```

Example: Australian population

```
{r holt-cmp-plot, echo=TRUE, dependson='holt-fit', fig.height=4.5}  
components(fit) |> autoplot()
```

Example: Australian population

```
{r holt-fc, echo=TRUE, dependson='holt-fit', fig.height=3.4} fit  
|>   forecast(h = 20) |>   autoplot(aus_economy) +   labs(y =  
"Population", x = "Year")
```

ETS(A,Ad,N): Damped trend method

Additive errors

Forecast equation

$$\hat{y}_{T+h|T} = \ell_T + (\phi + \dots + \phi^{h-1})b_T$$

Measurement equation

$$y_t = (\ell_{t-1} + \phi b_{t-1}) + \varepsilon_t$$

State equations

$$\ell_t = (\ell_{t-1} + \phi b_{t-1}) + \alpha \varepsilon_t$$

$$b_t = \phi b_{t-1} + \beta \varepsilon_t$$

ETS(A,Ad,N): Damped trend method

Additive errors

Forecast equation

$$\hat{y}_{T+h|T} = \ell_T + (\phi + \dots + \phi^{h-1})b_T$$

Measurement equation

$$y_t = (\ell_{t-1} + \phi b_{t-1}) + \varepsilon_t$$

State equations

$$\ell_t = (\ell_{t-1} + \phi b_{t-1}) + \alpha \varepsilon_t$$

$$b_t = \phi b_{t-1} + \beta \varepsilon_t$$

- Damping parameter $0 < \phi < 1$.
- If $\phi = 1$, identical to Holt's linear trend.
- As $h \rightarrow \infty$, $\hat{y}_{T+h|T} \rightarrow \ell_T + \phi b_T / (1 - \phi)$.
- Short-run forecasts trended, long-run forecasts constant.

Example: Australian population

```
{r, echo=TRUE, fig.height=3.4} aus_economy |> model(holt =  
ETS(Pop ~ trend("Ad")))) |> forecast(h = 20) |>  
autoplot(aus_economy)
```

Example: National populations

```
{r popfit, echo=TRUE} fit <- global_economy |> mutate(Pop = Population / 1e6) |>  
model(ets = ETS(Pop)) fit
```

Example: National populations

```
{r popfc, echo=TRUE, dependson="popfit"} fit |> forecast(h = 5)
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Lab Session 14

Try forecasting the Australian GDP from the `global_economy` data set using an ETS model.

Experiment with the various options in the `ETS()` function to see how much the forecasts change with damped trend, or with a Box-Cox transformation. Try to develop an intuition of what each is doing to the forecasts.

[Hint: use `h=20` when forecasting, so you can clearly see the differences between the various options when plotting the forecasts.]

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

ETS(A,A,A): Holt-Winters additive method

Forecast equation

$$\hat{y}_{t+h|t} = \ell_t + hb_t + s_{t+h-m(k+1)}$$

Observation equation

$$y_t = \ell_{t-1} + b_{t-1} + s_{t-m} + \varepsilon_t$$

State equations

$$\ell_t = \ell_{t-1} + b_{t-1} + \alpha \varepsilon_t$$

$$b_t = b_{t-1} + \beta \varepsilon_t$$

$$s_t = s_{t-m} + \gamma \varepsilon_t$$

- $k = \text{integer part of } (h - 1)/m$.
- $\sum_i s_i \approx 0$.
- Parameters: $0 \leq \alpha \leq 1$, $0 \leq \beta^* \leq 1$, $0 \leq \gamma \leq 1 - \alpha$ and $m = \text{period of seasonality}$ (e.g. $m = 4$ for quarterly data).

ETS(M,A,M): Holt-Winters multiplicative method

Forecast equation $\hat{y}_{t+h|t} = (\ell_t + hb_t)s_{t+h-m(k+1)}$

Observation equation $y_t = (\ell_{t-1} + b_{t-1})s_{t-m}(1 + \varepsilon_t)$

State equations $\ell_t = (\ell_{t-1} + b_{t-1})(1 + \alpha\varepsilon_t)$

$$b_t = b_{t-1}(1 + \beta\varepsilon_t)$$

$$s_t = s_{t-m}(1 + \gamma\varepsilon_t)$$

- k is integer part of $(h - 1)/m$.
- $\sum_i s_i \approx m$.
- Parameters: $0 \leq \alpha \leq 1$, $0 \leq \beta^* \leq 1$, $0 \leq \gamma \leq 1 - \alpha$ and m = period of seasonality (e.g. $m = 4$ for quarterly data).

Example: Australian holiday tourism

```
{r ausholidays-fit, echo=TRUE} holidays <- tourism |> filter(Purpose == "Holiday")  
fit <- holidays |> model(ets = ETS(Trips)) fit
```

Example: Australian holiday tourism

```
fit |>  
  filter(Region == "Snowy Mountains") |>  
  report()
```

Example: Australian holiday tourism

```
fit |>  
  filter(Region == "Snowy Mountains") |>  
components(fit)
```

Example: Australian holiday tourism

```
#| fig-height: 3.6
fit |>
  filter(Region == "Snowy Mountains") |>
  components(fit) |>
  autoplot()
```

Example: Australian holiday tourism

```
fit |> forecast()
```

Example: Australian holiday tourism

```
{r ausholidays-forecast-plot, fig.height=3.4} fit |> forecast() |> filter(Region == "Snowy Mountains") |> autoplot(holidays) + labs(x = "Year", y = "Overnight trips (thousands)")
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Exponential smoothing models

Additive Error

Trend Component

N (None)

A (Additive)

A_d (Additive damped)

Seasonal Component

	N (None)	A (Additive)	M (Multiplicative)
--	-------------	-----------------	-----------------------

A,N,N A,N,A ~~A,N,M~~

A,A,N A,A,A ~~A,A,M~~

A, A_d ,N A, A_d ,A ~~A, A_d ,M~~

Multiplicative Error

Trend Component

N (None)

A (Additive)

A_d (Additive damped)

Seasonal Component

	N (None)	A (Additive)	M (Multiplicative)
--	-------------	-----------------	-----------------------

M,N,N M,N,A M,N,M

M,A,N M,A,A M,A,M

M, A_d ,N M, A_d ,A M, A_d ,M

Estimating ETS models

- Smoothing parameters α, β, γ and ϕ , and the initial states $\ell_0, b_0, s_0, s_{-1}, \dots, s_{-m+1}$ are estimated by maximising the “likelihood” = the probability of the data arising from the specified model.
- For models with additive errors equivalent to minimising SSE.
- For models with multiplicative errors, **not** equivalent to minimising SSE.

Model selection

Akaike's Information Criterion

$$AIC = -2 \log(L) + 2k$$

where L is the likelihood and k is the number of parameters & initial states estimated in the model.

Model selection

Akaike's Information Criterion

$$AIC = -2 \log(L) + 2k$$

where L is the likelihood and k is the number of parameters & initial states estimated in the model.

Corrected AIC

$$AIC_c = AIC + \frac{2k(k + 1)}{T - k - 1}$$

which is the AIC corrected (for small sample bias).

Model selection

Akaike's Information Criterion

$$AIC = -2 \log(L) + 2k$$

where L is the likelihood and k is the number of parameters & initial states estimated in the model.

Corrected AIC

$$AIC_c = AIC + \frac{2k(k + 1)}{T - k - 1}$$

which is the AIC corrected (for small sample bias).

Bayesian Information Criterion

$$BIC = AIC + k(\log(T) - 2).$$

AIC and cross-validation

Minimizing the AIC assuming Gaussian residuals is asymptotically equivalent to minimizing one-step time series cross validation MSE.

Automatic forecasting

From Hyndman et al. (IJF, 2002):

- 1 Apply each model that is appropriate to the data.
Optimize parameters and initial values using MLE.
 - 2 Select best method using AICc.
 - 3 Produce forecasts using best method.
 - 4 Obtain forecast intervals using underlying state space model.
- Method performed very well in M3 competition.
 - Used as a benchmark in the M4 competition.

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Lab Session 15

Find an ETS model for the Gas data from aus_production.

- Why is multiplicative seasonality necessary here?
- Experiment with making the trend damped. Does it improve the forecasts?

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Non-Gaussian forecast distributions

```
{r cafe, fig.height=3.1} vic_cafe <- tsibbledata::aus_retail |> filter(State ==  
"Victoria", Industry == "Cafes, restaurants and catering services") |>  
select(Month, Turnover) vic_cafe |> autoplot(Turnover) + labs(title = "Monthly  
turnover of Victorian cafes")
```

Forecasting with transformations

```
vic_cafe |> autoplot(box_cox(Turnover, lambda = 0.2))
```

Forecasting with transformations

```
fit <- vic_cafe |>
  model(ets = ETS(box_cox(Turnover, 0.2)))
fit

{r include=FALSE} if (!identical(fabletools:::model_sum(fit$ets[[1]]), "ETS(A,A,A)")) {
  stop("Model not ETS(A,A,A)") }

(fc <- fit |> forecast(h = "3 years"))
```

Forecasting with transformations

```
fit <- vic_cafe |>
  model(ets = ETS(box_cox(Turnover, 0.2)))
fit

{r include=FALSE} if (!identical(fabletools:::model_sum(fit$ets[[1]]), "ETS(A,A,A)")) {
  stop("Model not ETS(A,A,A)") }
(fc <- fit |> forecast(h = "3 years"))
```

- $t(N)$ denotes a transformed normal distribution.
- back-transformation and bias adjustment is done automatically.

Forecasting with transformations

```
fc |> autoplot(vic_cafe)
```

Bootstrapped forecast distributions

```
sim <- fit |> generate(h = "3 years", times = 5, bootstrap = TRUE)  
sim
```

Bootstrapped forecast distributions

```
{r, fig.height=2.7} vic_cafe |> filter(year(Month) >= 2008) |> ggplot(aes(x = Month)) + geom_line(aes(y = Turnover)) + geom_line(aes(y = .sim, colour = as.factor(.rep)), data = sim) + labs(title = "Monthly turnover of Victorian cafes") + guides(col = FALSE)
```

Bootstrapped forecast distributions

```
fc <- fit |> forecast(h = "3 years", bootstrap = TRUE)  
fc
```

Bootstrapped forecast distributions

```
fc |> autoplot(vic_cafe) +  
  labs(title = "Monthly turnover of Victorian cafes")
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL Features
- Lab Session 9
- Dimension reduction for features

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

ARIMA models

- AR:** autoregressive (lagged observations as inputs)
- I:** integrated (differencing to make series stationary)
- MA:** moving average (lagged errors as inputs)

ARIMA models

- AR:** autoregressive (lagged observations as inputs)
- I:** integrated (differencing to make series stationary)
- MA:** moving average (lagged errors as inputs)

An ARIMA model is rarely interpretable in terms of visible data structures like trend and seasonality. But it can capture a huge range of time series patterns.

Stationarity

Definition

If $\{y_t\}$ is a stationary time series, then for all s , the distribution of (y_t, \dots, y_{t+s}) does not depend on t .

Stationarity

Definition

If $\{y_t\}$ is a stationary time series, then for all s , the distribution of (y_t, \dots, y_{t+s}) does not depend on t .

A **stationary series** is:

- roughly horizontal
- constant variance
- no patterns predictable in the long-term

Stationary?

```
{r, fig.height=2.7} gafa_stock |> filter(Symbol == "GOOG", year(Date) == 2018) |> autoplot(Close) + labs(y = "Google closing stock price ($US)")
```

Stationary?

```
{r, fig.height=2.7} gafa_stock |> filter(Symbol == "GOOG", year(Date) == 2018) |> autoplot(difference(Close)) + labs(y = "Daily change in Google closing stock price")
```

Differencing

- Differencing helps to **stabilize the mean**.
- The differenced series is the *change* between each observation in the original series.
- Occasionally the differenced data will not appear stationary and it may be necessary to difference the data a second time.
- In practice, it is almost never necessary to go beyond second-order differences.

Autoregressive models

Autoregressive (AR) models:

$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_p y_{t-p} + \varepsilon_t,$$

where ε_t is white noise. A multiple regression with **lagged values** of y_t as predictors.

```
{r arp, echo=FALSE, fig.height=2.3} set.seed(1) p1 <-  
tsibble(idx = seq_len(100), sim = 10 +  
arima.sim(list(ar = -0.8), n = 100), index = idx) |>  
autoplot(sim) + labs(x = "time", y = "", title =  
"AR(1)") p2 <- tsibble(idx = seq_len(100), sim = 20 +  
arima.sim(list(ar = c(1.3, -0.7)), n = 100), index =
```

Moving Average (MA) models

Moving Average (MA) models:

$$y_t = c + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \cdots + \theta_q \varepsilon_{t-q},$$

where ε_t is white noise. A multiple regression with **lagged errors** as predictors. *Don't confuse with moving average smoothing!*

```
“{r maq, fig.height=2.3, echo=FALSE} set.seed(2) p1 <-  
tsibble(idx = seq_len(100), sim = 20 + arima.sim(list(ma = 0.8),  
n = 100), index = idx) |> autoplot(sim) + labs(x = “time”, y = “”,  
title =”MA(1)“) p2 <- tsibble(idx = seq_len(100), sim =  
arima.sim(list(ma = c(-1, +0.8)), n = 100), index = idx) |>
```

Example: National populations

```
{r popfit3, echo=TRUE} fit |> filter(Country == "Australia") |>  
report()
```

Example: National populations

```
{r popfit3, echo=TRUE} fit |> filter(Country == "Australia") |>  
report()
```

$$y_t = 2y_{t-1} - y_{t-2} - 0.7\varepsilon_{t-1} + \varepsilon_t$$
$$\varepsilon_t \sim \text{NID}(0, 4 \times 10^9)$$

Understanding ARIMA models

- If $c = 0$ and $d = 0$, the long-term forecasts will go to zero.
- If $c = 0$ and $d = 1$, the long-term forecasts will go to a non-zero constant.
- If $c = 0$ and $d = 2$, the long-term forecasts will follow a straight line.
- If $c \neq 0$ and $d = 0$, the long-term forecasts will go to the mean of the data.
- If $c \neq 0$ and $d = 1$, the long-term forecasts will follow a straight line.
- If $c \neq 0$ and $d = 2$, the long-term forecasts will follow a quadratic trend.

Understanding ARIMA models

Forecast variance and d

- The higher the value of d , the more rapidly the prediction intervals increase in size.
- For $d = 0$, the long-term forecast standard deviation will go to the standard deviation of the historical data.

Example: National populations

```
{r popfc2, echo=TRUE, fig.height=3.4} fit |> forecast(h = 10) |> filter(Country == "Australia") |> autoplot(global_economy)
```

How does ARIMA() work?

Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d via KPSS test.
- Select p , q and inclusion of c by minimising AICc.
- Use stepwise search to traverse model space.

How does ARIMA() work?

Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d via KPSS test.
- Select p, q and inclusion of c by minimising AICc.
- Use stepwise search to traverse model space.

$$\text{AICc} = -2 \log(L) + 2(p + q + k + 1) \left[1 + \frac{(p + q + k + 2)}{T - p - q - k - 2} \right]$$

where L is the maximised likelihood fitted to the *differenced* data, $k = 1$ if $c \neq 0$ and $k = 0$ otherwise.

How does ARIMA() work?

Hyndman and Khandakar (JSS, 2008) algorithm:

- Select no. differences d via KPSS test.
- Select p, q and inclusion of c by minimising AICc.
- Use stepwise search to traverse model space.

$$\text{AICc} = -2 \log(L) + 2(p + q + k + 1) \left[1 + \frac{(p + q + k + 2)}{T - p - q - k - 2} \right]$$

where L is the maximised likelihood fitted to the *differenced* data, $k = 1$ if $c \neq 0$ and $k = 0$ otherwise.

Note: Can't compare AICc for different values of d .

How does ARIMA() work?

Step1: Select current model (with smallest AICc) from:

ARIMA(2, d, 2)

ARIMA(0, d, 0)

ARIMA(1, d, 0)

ARIMA(0, d, 1)

How does ARIMA() work?

Step1: Select current model (with smallest AICc) from:

ARIMA(2, d, 2)

ARIMA(0, d, 0)

ARIMA(1, d, 0)

ARIMA(0, d, 1)

Step 2: Consider variations of current model:

- vary one of p, q , from current model by ± 1 ;
- p, q both vary from current model by ± 1 ;
- Include/exclude c from current model.

Model with lowest AICc becomes current model.

How does ARIMA() work?

Step1: Select current model (with smallest AICc) from:

ARIMA(2, d, 2)

ARIMA(0, d, 0)

ARIMA(1, d, 0)

ARIMA(0, d, 1)

Step 2: Consider variations of current model:

- vary one of p, q , from current model by ± 1 ;
- p, q both vary from current model by ± 1 ;
- Include/exclude c from current model.

Model with lowest AICc becomes current model.

Repeat Step 2 until no lower AICc can be found.

How does ARIMA() work?

```
{r ARMAgridsearch, echo=FALSE, message=FALSE,
warning=FALSE, fig.asp=1, out.width="60%",
fig.width=4, fig.height=4} start <- tribble( ~p, ~q,
0, 0, 1, 0, 0, 1, 2, 2 ) selected <- tribble(
~p, ~q, 2, 2, 3, 3, 4, 2 ) griddf <-
expand.grid(p = 0:6, q = 0:6) |> as_tibble() |>
left_join(start |> mutate(start = TRUE)) |>
left_join(selected |> mutate(chosen = TRUE)) |>
replace_na(list(start = FALSE, chosen = FALSE)) |>
mutate( step = case_when( start ~ 1,
(p - selected$p[1])^2 + (q - selected$q[1])^2 <= 2 ~
```

How does ARIMA() work?

```
{r ARMAgridsearch2, echo=FALSE, message=FALSE,  
warning=FALSE, fig.asp=1, out.width="60%",  
fig.width=4, fig.height=4} griddf |> ggplot(aes(x =  
q, y = p)) + geom_point(aes(alpha = 0.2), colour =  
"gray", size = 5, shape = 19) + geom_segment(aes(x =  
fromq, y = fromp, xend = q, yend = p, col = step),  
data = griddf |> filter(step == "2"), arrow =  
arrow(length = unit(0.15, "inches"), type = "open"),  
size = 1, lineend = "butt") + geom_point(aes(col  
= step), size = 5, shape = 19, data = griddf  
|> filter(step %in% c("1", "2")) ) +
```

How does ARIMA() work?

```
{r ARMAgridsearch3, echo=FALSE, message=FALSE,  
warning=FALSE, fig.asp=1, out.width="60%",  
fig.width=4, fig.height=4} griddf |> ggplot(aes(x =  
q, y = p)) + geom_point(aes(alpha = 0.2), colour =  
"gray", size = 5, shape = 19) + geom_segment(aes(x =  
fromq, y = fromp, xend = q, yend = p, col = step),  
data = griddf |> filter(step %in% "3"), arrow =  
arrow(length = unit(0.15, "inches"), type = "open"),  
size = 1, lineend = "butt") + geom_point(aes(col  
= step), size = 5, shape = 19, data = griddf  
|> filter(step %in% c("1", "2", "3")) ) +
```

How does ARIMA() work?

```
{r ARMAgridsearch4, echo=FALSE, message=FALSE,  
warning=FALSE, fig.asp=1, out.width="60%",  
fig.width=4, fig.height=4} griddf |> ggplot(aes(x =  
q, y = p)) + geom_point(aes(alpha = 0.2), colour =  
"gray", size = 5, shape = 19) + geom_segment(aes(x =  
fromq, y = fromp, xend = q, yend = p, col = step),  
data = griddf |> filter(step %in% "4"), arrow =  
arrow(length = unit(0.15, "inches"), type = "open"),  
size = 1, lineend = "butt") + geom_point(aes(col  
= step), size = 5, shape = 19, data = griddf  
|> filter(step %in% c("1", "2", "3", "4")) ) +
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL Features
- Lab Session 9
- Dimension reduction for features

Lab Session 16

For the Australian job vacancies:

```
{r, eval = FALSE} skill_vacancies <-
readxl::read_excel("data/Internet Vacancies, ANZSCO
Skill Level, States and Territories - May 2024.xlsx",
sheet = 2) |>      # Tidy into a long form
pivot_longer(matches("\d{5}"), names_to = "month",
values_to = "vacancies",
names_transform = list(month = ~
yearmonth(as.Date(as.integer(.)), origin =
"1900-01-01")))) |>      # Remove aggregate vacancies
filter(State != "AUST", Skill level != 0) |>      #
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Seasonal ARIMA models

ARIMA	$\underbrace{(p, d, q)}_{\substack{\uparrow \\ \text{Non-seasonal part} \\ \text{of the model}}}$	$\underbrace{(P, D, Q)_m}_{\substack{\uparrow \\ \text{Seasonal part of} \\ \text{of the model}}}$
-------	---	--

- m = number of observations per year.
- d first differences, D seasonal differences
- p AR lags, q MA lags
- P seasonal AR lags, Q seasonal MA lags

Seasonal and non-seasonal terms combine multiplicatively

Cortecosteroid drug sales

```
{r, echo=TRUE} h02 <- PBS |> filter(ATC2 == "H02")  
|> summarise(Cost = sum(Cost) / 1e6)
```

Cortecosteroid drug sales

```
{r, echo=TRUE} h02 |> autoplot( Cost )
```

Cortecosteroid drug sales

```
{r, echo=TRUE} h02 |> autoplot(log(Cost))
```

Cortecosteroid drug sales

```
{r, echo=TRUE} h02 |> autoplot( log(Cost) |>  
difference(12) )
```

Cortecosteroid drug sales

```
{r, echo=TRUE} h02 |> autoplot( log(Cost) |>  
difference(12) |> difference(1) )
```

Example: US electricity production

```
{r h02fit, echo=TRUE, fig.height=2.8} h02 |> model(arima =  
ARIMA(log(Cost))) |> report()
```

Example: US electricity production

```
{r h02fcst, echo=TRUE, fig.height=2.8} h02 |> model(arima =  
ARIMA(log(Cost))) |> forecast(h = "3 years") |> autoplot(h02)
```

Cortecosteroid drug sales

```
{r h02tryharder, echo=TRUE, fig.height=3.6} fit <- h02 |> model(best =  
ARIMA(log(Cost),      stepwise = FALSE,      approximation = FALSE,      order_constraint  
= p + q + P + Q <= 9    )) report(fit)
```

Cortecosteroid drug sales

```
{r h02f, echo=TRUE, fig.height=2.8} fit |> forecast() |>  
autoplot(h02) + labs(y = "H02 Expenditure ($AUD)", x = "Year")
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL Features
- Lab Session 9
- Dimension reduction for features

Lab Session 17

For the Australian tourism data (from `tourism`):

- Fit a suitable ARIMA model for all data.
- Produce forecasts of your fitted models.
- Check the forecasts for the “Snowy Mountains” and “Melbourne” regions. Do they look reasonable?

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Forecast ensembles

```
{r trainall, echo=TRUE, dependson='tourism'} train <- tourism |>  
filter(year(Quarter) <= 2014) fit <- train |>  model(      ets = ETS(Trips),  
arima = ARIMA(Trips),      snaive = SNAIVE(Trips)    ) |>  mutate(mixed = (ets  
+ arima + snaive) / 3)
```

- Ensemble forecast `mixed` is a simple average of the three fitted models.
- `forecast()` will produce distributional forecasts taking into account the correlations between the forecast errors of the component models.

Forecast ensembles

```
{r trainfc, dependson='trainall'} #| fig-height: 3.6 fc <- fit |> forecast(h =  
"3 years") fc |> filter(Region == "Snowy Mountains", Purpose == "Holiday")  
|> autoplot(tourism, level = NULL)
```

Forecast ensembles

```
{r snowy-test-accuracy, dependson='trainfc'} accuracy(fc, tourism) |>  
group_by(.model) |> summarise(      RMSE = mean(RMSE),      MAE = mean(MAE),  
MASE = mean(MASE) ) |> arrange(RMSE)
```

Forecast ensembles

Can we do better than equal weights?

Forecast ensembles

Can we do better than equal weights?

- Hard to find weights that improve forecast accuracy.
- Known as the “forecast combination puzzle”.
- Solution: FFOMA

Forecast ensembles

Can we do better than equal weights?

- Hard to find weights that improve forecast accuracy.
- Known as the “forecast combination puzzle”.
- Solution: FFORMA

FFORMA (Feature-based FOREcast Model Averaging)

- Vector of time series features used to predict best weights.
- A modification of xgboost is used.
- Method came 2nd in the 2018 M4 international forecasting competition.
- Main author: Pablo Montero-Manso (now Uni Sydney)
- Not (yet) available for fable.

Outline

- 1 Time series data and tsibbles
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Lab Session 1
- 5 Time plots
- 6 Lab Session 2
- 7 Seasonal plots
- 8 Lab Session 3
- 9 Seasonal or cyclic?
- 10 Lag plots and autocorrelation
- 11 Lab Session 4
- 12 White noise
- 13 Lab Session 5
- 14 Per capita adjustments
- 15 Lab Session 6
- 16 Inflation adjustments
- 17 Mathematical transformations
- 18 Lab Session 7
- 19 Time series decompositions
- 20 Lab Session 8
- 21 Multiple seasonality
- 22 A cautionary tale
- 23 Notice: Material planned to change
- 24 STL features
- 25 Lab Session 9
- 26 Dimension reduction for features

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Notice: Material planned to change

This material is planned to be updated to better align with the training needs of the Department of Education.

In particular, the new material will be more focused on:

- the use of policy in models,
- forecasting with different scenarios.

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Regression with ARIMA errors

Regression models

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \varepsilon_t,$$

- y_t modeled as function of k explanatory variables
- In regression, we assume that ε_t is white noise.

Regression with ARIMA errors

Regression models

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \varepsilon_t,$$

- y_t modeled as function of k explanatory variables
- In regression, we assume that ε_t is white noise.

RegARIMA model

$$y_t = \beta_0 + \beta_1 x_{1,t} + \cdots + \beta_k x_{k,t} + \eta_t,$$

$$\eta_t \sim \text{ARIMA}$$

- Residuals are from ARIMA model.
- Estimate model in one step using MLE
- Select model with lowest AICc value.

US personal consumption and income

```
{r usconsumpdata, echo=TRUE} us_change
```

US personal consumption and income

```
{r usconsump, fig.height=4, fig.width=6,  
out.height="55%"} us_change |>  
pivot_longer(-Quarter, names_to = "variable",  
values_to = "value") |> ggplot(aes(y = value, x =  
Quarter, group = variable)) + geom_line() +  
facet_grid(variable ~ ., scales = "free_y") + labs(x  
= "Year", y = "", title = "Quarterly changes in  
US consumption and personal income")
```

US personal consumption and income

```
{r usconsump_pairs, fig.height=6, fig.width=8.5, echo=TRUE} #| out-width: 70%
us_change |> as_tibble() |> select(-Quarter) |> GGally::ggpairs()
```

US personal consumption and income

- No need for transformations or further differencing.
- Increase in income does not necessarily translate into instant increase in consumption (e.g., after the loss of a job, it may take a few months for expenses to be reduced to allow for the new circumstances). We will ignore this for now.

US personal consumption and income

```
{r usconsump2, echo=TRUE, fig.height=3} fit <- us_change |> model(regarima =  
ARIMA(Consumption ~ Income + Production + Savings +  
Unemployment)) report(fit)
```

US personal consumption and income

```
{r usconsump2, echo=TRUE, fig.height=3} fit <- us_change |> model(regarima =  
ARIMA(Consumption ~ Income + Production + Savings +  
Unemployment)) report(fit)
```

Write down the equations for the fitted model.

US personal consumption and income

```
{r , echo=TRUE, fig.height=3.7,  
dependson='usconsump2'} gg_tsresiduals(fit)
```

US personal consumption and income

```
{r , echo=TRUE, fig.height=3.7, dependson='usconsump2'} augment(fit) |>  
features(.resid, ljung_box, dof = 2, lag = 12)
```

US personal consumption and income

```
{r usconsump3, echo=TRUE, fig.height=2.4,
dependson='usconsump2'} us_change_future <-
new_data(us_change, 8) |>  mutate(Income =
tail(us_change$Income, 1),           Production =
tail(us_change$Production, 1),       Savings =
tail(us_change$Savings, 1),          Unemployment =
tail(us_change$Unemployment, 1)) forecast(fit,
new_data = us_change_future) |>  autoplot(us_change)
+  labs(x = "Year", y = "Percentage change",
title = "Forecasts from dynamic regression")
```

Forecasting

- To forecast a regression model with ARIMA errors, we need to forecast the regression part of the model and the ARIMA part of the model and combine the results.
- Some predictors are known into the future (e.g., time, dummies).
- Separate forecasting models may be needed for other predictors.
- Forecast intervals ignore the uncertainty in forecasting the predictors.

Daily electricity demand

Model daily electricity demand as a function of temperature using quadratic regression with ARMA errors.

```
{r, echo=TRUE, fig.height=2.7} vic_elec_daily |> ggplot(aes(x = Temperature, y = Demand, colour = Day_Type)) + geom_point() + labs(x = "Maximum temperature", y = "Electricity demand (GW)")
```

Daily electricity demand

```
{r, echo=TRUE, fig.height=3.2} vic_elec_daily |> pivot_longer(c(Demand, Temperature)) |>  
ggplot(aes(x = Date, y = value)) + geom_line() + facet_grid(vars(name), scales = "free_y")
```

Daily electricity demand

```
{r dailymodel, echo=TRUE} fit <- vic_elec_daily |> model(fit = ARIMA(Demand ~ Temperature +  
I(Temperature^2) + (Day_Type == "Weekday"))) report(fit)
```

Daily electricity demand

```
{r, echo=TRUE, dependson='dailymodel'} augment(fit) |>  
gg_tsdisplay(.resid, plot_type = "histogram")
```

Daily electricity demand

```
{r, echo=TRUE, dependson='dailymodel'} augment(fit) |>  
features(.resid, ljung_box, dof = 9, lag = 14)
```

Daily electricity demand

```
{r, echo=TRUE, dependson='dailymodel'} # Forecast one day ahead  
vic_next_day <- new_data(vic_elec_daily, 1) |> mutate(Temperature =  
26, Day_Type = "Holiday") forecast(fit, vic_next_day)
```

Daily electricity demand

```
{r, echo=TRUE} vic_elec_future <-
new_data(vic_elec_daily, 14) |> mutate(
Temperature = 26, Holiday = c(TRUE, rep(FALSE,
13)), Day_Type = case_when( Holiday ~
"Holiday", wday(Date) %in% 2:6 ~ "Weekday",
TRUE ~ "Weekend" ) )
```

Daily electricity demand

```
{r, echo = TRUE, dependson='dailymodel'} forecast(fit, vic_elec_future) |>  
autoplot(vic_elec_daily) + labs(y = "Electricity demand (GW)")
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Lab Session 18

Repeat the daily electricity example, but instead of using a quadratic function of temperature, use a piecewise linear function with the “knot” around 20 degrees Celsius (use predictors Temperature & Temp2). How can you optimize the choice of knot?

```
{r echo=TRUE, eval=FALSE} vic_elec_daily <- vic_elec |> filter(year(Time) == 2014) |>  
index_by(Date = date(Time)) |> summarise(Demand = sum(Demand) / 1e3, Temperature =  
max(Temperature), Holiday = any(Holiday) ) |> mutate(Temp2 = I(pmax(Temperature -  
20, 0)), Day_Type = case_when( Holiday ~ "Holiday", wday(Date)  
%in% 2:6 ~ "Weekday", TRUE ~ "Weekend") )
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL Features
- Lab Session 9
- Dimension reduction for features

Dynamic harmonic regression

Combine Fourier terms with ARIMA errors

Advantages

- it allows any length seasonality;
- for data with more than one seasonal period, you can include Fourier terms of different frequencies;
- the seasonal pattern is smooth for small values of K (but more wiggly seasonality can be handled by increasing K);
- the short-term dynamics are easily handled with a simple ARMA error.

Disadvantages

- seasonality is assumed to be fixed

Eating-out expenditure

```
{r cafe, echo=TRUE, fig.height=2.3, fig.width=8} aus_cafe <- aus_retail |> filter(Industry == "Cafes, restaurants and takeaway food services", year(Month) %in% 2004:2018 ) |> summarise(Turnover = sum(Turnover)) aus_cafe |> autoplot(Turnover)
```

Eating-out expenditure

```
{r cafefit, dependson='cafe', fig.height=5, echo=TRUE, results='hide'} fit <- aus_cafe  
|> model(`K = 1` = ARIMA(log(Turnover) ~ fourier(K = 1) + PDQ(0, 0, 0)), `K = 2`  
= ARIMA(log(Turnover) ~ fourier(K = 2) + PDQ(0, 0, 0)), `K = 3` =  
ARIMA(log(Turnover) ~ fourier(K = 3) + PDQ(0, 0, 0)), `K = 4` = ARIMA(log(Turnover)  
~ fourier(K = 4) + PDQ(0, 0, 0)), `K = 5` = ARIMA(log(Turnover) ~ fourier(K = 5) +  
PDQ(0, 0, 0)), `K = 6` = ARIMA(log(Turnover) ~ fourier(K = 6) + PDQ(0, 0, 0)) )  
glance(fit) {r, echo = FALSE} glance(fit) |> select(.model, sigma2, log_lik, AIC,  
AICc, BIC) |> knitr::kable()
```

Eating-out expenditure

```
{r, include=FALSE} cafe_plot <- function(...) {    fit  
|>      select(...) |>      forecast() |>  
autoplot(aus_cafe) +      labs(title = sprintf("Log  
transformed %s, fourier(K = %s)",  
model_sum(select(fit, ...)[[1]][[1]]), deparse(..1)))  
+      geom_label(      aes(x = yearmonth("2007 Jan"),  
y = 4250, label = paste0("AICc = ", format(AICc))),  
data = glance(select(fit, ...)) ) +  
geom_line(aes(y = .fitted), colour = "red",  
augment(select(fit, ...))) +      ylim(c(1500, 5100)) }  
{r cafe1, dependson='cafe', fig.height=5, echo=FALSE}
```

Eating-out expenditure

```
{r cafe2, dependson='cafe', fig.height=5, echo=FALSE}  
cafe_plot("K = 2")
```

Eating-out expenditure

```
{r cafe3, dependson='cafe', fig.height=5, echo=FALSE}  
cafe_plot("K = 3")
```

Eating-out expenditure

```
{r cafe4, dependson='cafe', fig.height=5, echo=FALSE}  
cafe_plot("K = 4")
```

Eating-out expenditure

```
{r cafe5, dependson='cafe', fig.height=5, echo=FALSE}  
cafe_plot("K = 5")
```

Eating-out expenditure

```
{r cafe6, dependson='cafe', fig.height=5, echo=FALSE}  
cafe_plot("K = 6")
```

Example: weekly gasoline products

```
{r, echo = FALSE} options(width = 70)

{r gasmodel, echo=TRUE} fit <- us_gasoline |> model(ARIMA(Barrels ~ fourier(K = 13) + PDQ(0, 0,
0))) report(fit)
```

Example: weekly gasoline products

```
{r gasf, echo=TRUE, fig.height=3} forecast(fit, h = "3 years") |>  
autoplot(us_gasoline)
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Lab Session 19

Repeat Lab Session 18 but using all available data, and handling the annual seasonality using Fourier terms.

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Lagged predictors

Sometimes a change in x_t does not affect y_t instantaneously

- y_t = sales, x_t = advertising.
- y_t = stream flow, x_t = rainfall.
- y_t = size of herd, x_t = breeding stock.

Lagged predictors

Sometimes a change in x_t does not affect y_t instantaneously

- y_t = sales, x_t = advertising.
 - y_t = stream flow, x_t = rainfall.
 - y_t = size of herd, x_t = breeding stock.
-
- These are dynamic systems with input (x_t) and output (y_t).
 - x_t is often a leading indicator.
 - There can be multiple predictors.

Lagged predictors

The model include present and past values of predictor:

$X_t, X_{t-1}, X_{t-2}, \dots$

$$y_t = a + \nu_0 X_t + \nu_1 X_{t-1} + \dots + \nu_k X_{t-k} + \eta_t$$

where η_t is an ARIMA process.

Lagged predictors

The model include present and past values of predictor:

$X_t, X_{t-1}, X_{t-2}, \dots$

$$y_t = a + \nu_0 X_t + \nu_1 X_{t-1} + \dots + \nu_k X_{t-k} + \eta_t$$

where η_t is an ARIMA process.

- x can influence y , but y is not allowed to influence x .

Example: Insurance quotes and TV adverts

insurance

Example: Insurance quotes and TV adverts

```
{r tvadvert, dependson='tvadvertdata', echo=FALSE}
insurance |>  pivot_longer(c(Quotes, TVadverts)) |>
ggplot(aes(x = Month, y = value)) +  geom_line() +
facet_grid(vars(name), scales = "free_y") +  labs(x =
"Year", y = NULL, title = "Insurance advertising and
quotations")
```

Example: Insurance quotes and TV adverts

```
{r tvadvertpairs, dependson='tvadvertdata',
echo=FALSE} insurance |>   mutate(    lag1 =
lag(TVadverts),    lag2 = lag(lag1)    ) |>
as_tibble() |>   select(-Month) |>   rename(lag0 =
TVadverts) |>   pivot_longer(-Quotes, names_to =
"Lag", values_to = "TV_advert") |>   ggplot(aes(x =
TV_advert, y = Quotes)) +   geom_point() +
facet_grid(. ~ Lag) +   labs(title = "Insurance
advertising and quotations")
```

Example: Insurance quotes and TV adverts

```
{r, echo=TRUE} fit <- insurance |> # Restrict data so models use same  
fitting period  mutate(Quotes = c(NA, NA, NA, Quotes[4:40])) |> model(  
ARIMA(Quotes ~ pdq(d = 0) + TVadverts),      ARIMA(Quotes ~ pdq(d = 0) +  
TVadverts +      lag(TVadverts)),      ARIMA(Quotes ~ pdq(d = 0) + TVadverts +  
lag(TVadverts) +      lag(TVadverts, 2)),      ARIMA(Quotes ~ pdq(d = 0) +  
TVadverts +      lag(TVadverts) +      lag(TVadverts, 2) +  
lag(TVadverts, 3))    )
```

Example: Insurance quotes and TV adverts

```
{r, echo=TRUE, results = 'hide'} glance(fit) {r, echo = FALSE} glance(fit) |>  
transmute(`Lag order` = 0:3, sigma2, log_lik, AIC, AICc, BIC) |>  
knitr::kable()
```

Example: Insurance quotes and TV adverts

```
{r tvadvertagain, echo=TRUE} # Re-fit to all data fit <- insurance |> model(ARIMA(Quotes ~ TVadverts + lag(TVadverts) + pdq(d = 0))) report(fit)
```

Example: Insurance quotes and TV adverts

```
{r tvadvertagain, echo=TRUE} # Re-fit to all data fit <- insurance |> model(ARIMA(Quotes ~ TVadverts + lag(TVadverts) + pdq(d = 0))) report(fit)

{r tvadvertparam, echo=FALSE, dependson='tvadvertagain'} # Store coefficients
coef <- rlang::set_names(tidy(fit)$estimate, tidy(fit)$term)
phi1 <- coef["ar1"] theta1 <- coef["ma1"]
theta2 <- coef["ma2"] intercept <- coef["intercept"]
gamma0 <- coef["TVadverts"] gamma1 <-
coef["lag(TVadverts)"]
```

$y_t = \text{rformat(intercept, digits = 3)} + \text{rformat(gamma0, digits = 3)}x_t + \text{rformat(phi1, digits = 3)}\eta_{t-1} + \varepsilon_t + \text{rformat(theta1, digits = 2)}\varepsilon_{t-1} + \text{rformat(theta2, digits = 2)}\varepsilon_{t-2}$

Example: Insurance quotes and TV adverts

```
{r, echo=TRUE, fig.height=3} advert_a <- new_data(insurance, 20) |>  
mutate(TVadverts = 10) forecast(fit, advert_a) |> autoplot(insurance)
```

Example: Insurance quotes and TV adverts

```
{r, echo=TRUE, fig.height=3} advert_b <- new_data(insurance, 20) |>  
mutate(TVadverts = 8) forecast(fit, advert_b) |> autoplot(insurance)
```

Example: Insurance quotes and TV adverts

```
{r, echo=TRUE, fig.height=3} advert_c <- new_data(insurance, 20) |>  
mutate(TVadverts = 6) forecast(fit, advert_c) |> autoplot(insurance)
```

Outline

- 1 Time series data and tsibbles
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Lab Session 1
- 5 Time plots
- 6 Lab Session 2
- 7 Seasonal plots
- 8 Lab Session 3
- 9 Seasonal or cyclic?
- 10 Lag plots and autocorrelation
- 11 Lab Session 4
- 12 White noise
- 13 Lab Session 5
- 14 Per capita adjustments
- 15 Lab Session 6
- 16 Inflation adjustments
- 17 Mathematical transformations
- 18 Lab Session 7
- 19 Time series decompositions
- 20 Lab Session 8
- 21 Multiple seasonality
- 22 A cautionary tale
- 23 Notice: Material planned to change
- 24 STL Features
- 25 Lab Session 9
- 26 Dimension reduction for features

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Notice: Material planned to change

This material is planned to be updated to better align with the training needs of the Department of Education.

In particular, this section will be reduced or removed to make time for multivariate econometric models.

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Australian Pharmaceutical Benefits Scheme



PBS sales

```
{r setwidth, include=FALSE} fred <- options(width = 71)  
{r pbs, dependson='setwidth'} PBS  
  
{r setwidthback, include=FALSE} options(width = fred$width)
```

ATC drug classification

- A Alimentary tract and metabolism
- B Blood and blood forming organs
- C Cardiovascular system
- D Dermatologicals
- G Genito-urinary system and sex hormones
- H Systemic hormonal preparations, excluding sex hormones and insulins
- J Anti-infectives for systemic use
- L Antineoplastic and immunomodulating agents
- M Musculo-skeletal system
- N Nervous system
- P Antiparasitic products, insecticides and repellents
- R Respiratory system
- S Sensory organs
- V Various

ATC drug classification

ATC1: 14 classes

A

Alimentary tract and metabolism

ATC2: 84 classes

A10

Drugs used in diabetes

A10B

Blood glucose lowering drugs

A10BA

Biguanides

A10BA02

Metformin

Australian tourism

```
#| label: ausmap
#| fig-height: 3
#| echo: false
library(sf)
# Use Okabe-Ito color-blind friendly color palette
state_colors <- c(
  `New South Wales` = "#56b4e9",
  `Victoria` = "#0072b2",
  `Queensland` = "#009e73",
  `South Australia` = "#f0e442",
  `Northern Territory` = "#d55e00",
  `Western Australia` = "#e69f00",
  `Tasmania` = "#cc79a7",
  `Australian Capital Territory` = "#cccccc"
)
read_sf("tourism/Tourism_Regions_2020.shp") |>
  rename(State = "STE_NAME16") |>
  ggplot() +
  geom_sf(aes(fill = State), alpha = 0.8) +
```

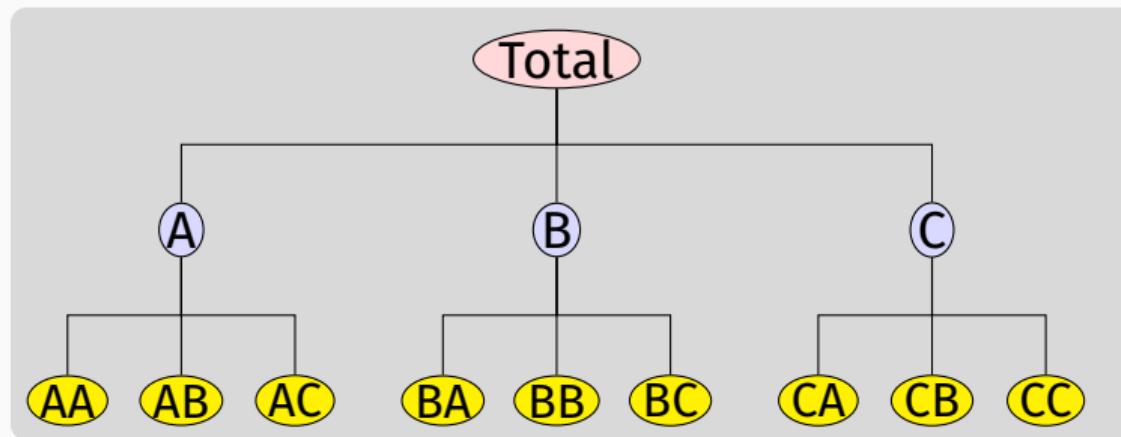
Australian tourism

tourism

- Quarterly data on visitor nights, 1998:Q1 – 2017:Q4
- From: *National Visitor Survey*, based on annual interviews of 120,000 Australians aged 15+, collected by Tourism Research Australia.
- Split by 8 states and 76 regions
- Split by purpose of travel
 - ▶ Holiday
 - ▶ Visiting friends and relatives (VFR)
 - ▶ Business
 - ▶ Other
- 304 bottom-level series

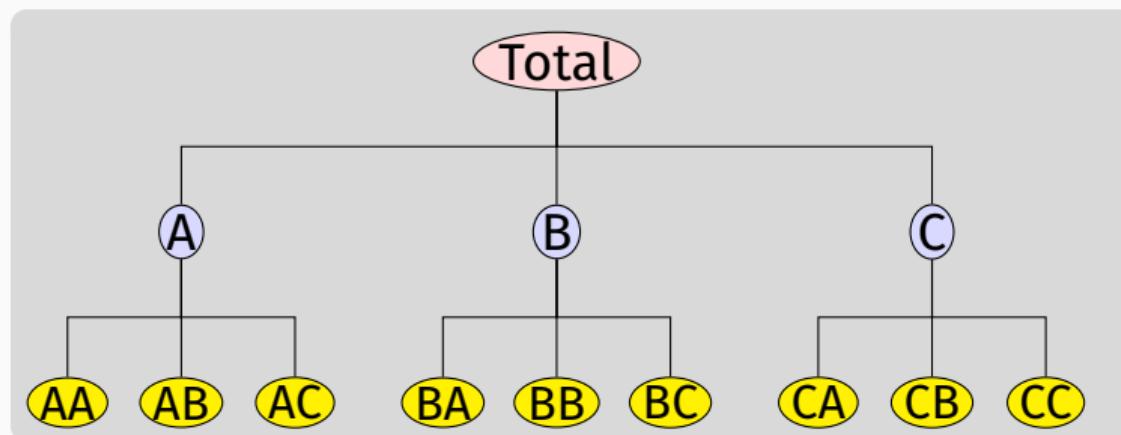
Hierarchical time series

A **hierarchical time series** is a collection of several time series that are linked together in a hierarchical structure.



Hierarchical time series

A **hierarchical time series** is a collection of several time series that are linked together in a hierarchical structure.

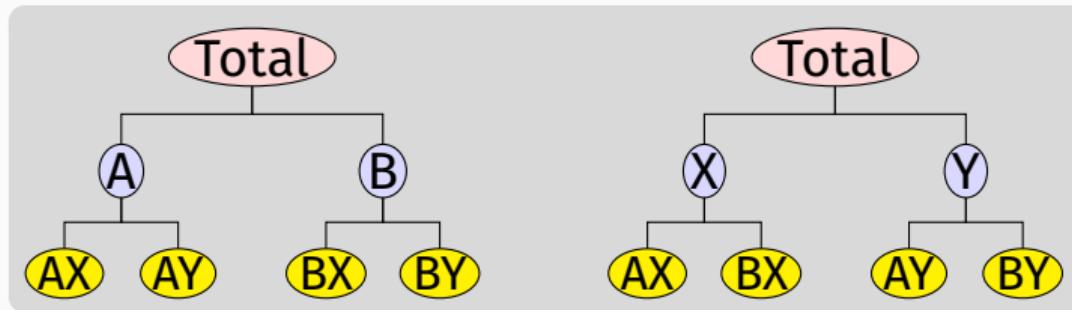


Examples

- PBS sales by ATC groups
- Tourism demand by states, regions

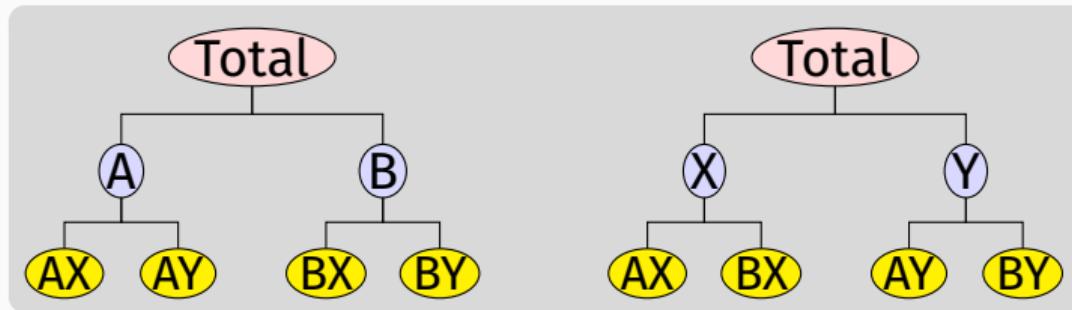
Grouped time series

A **grouped time series** is a collection of time series that can be grouped together in a number of non-hierarchical ways.



Grouped time series

A **grouped time series** is a collection of time series that can be grouped together in a number of non-hierarchical ways.



Examples

- Tourism by state and purpose of travel
- Retail sales by product groups/sub groups, and by countries/regions

Creating aggregates

```
PBS |>
  aggregate_key(ATC1 / ATC2, Scripts = sum(Scripts)) |>
  filter(Month == yearmonth("1991 Jul")) |>
  print(n = 18)
```

Creating aggregates

```
tourism |>
  aggregate_key(Purpose * (State / Region), Trips = sum(Trips)) |>
  filter(Quarter == yearquarter("1998 Q1")) |>
  print(n = 15)
```

Creating aggregates

- Similar to `summarise()` but using the key structure
- A grouped structure is specified using `grp1 * grp2`
- A nested structure is specified via parent / child.
- Groups and nesting can be mixed:
`(country/region/city) * (brand/product)`
- All possible aggregates are produced.
- These are useful when forecasting at different levels of aggregation.

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

The problem

- 1 How to forecast time series at all nodes such that the forecasts add up in the same way as the original data?
- 2 Can we exploit relationships between the series to improve the forecasts?

The problem

- 1 How to forecast time series at all nodes such that the forecasts add up in the same way as the original data?
- 2 Can we exploit relationships between the series to improve the forecasts?

The solution

- 1 Forecast all series at all levels of aggregation using an automatic forecasting algorithm.
(e.g., ETS, ARIMA, ...)
- 2 Reconcile the resulting forecasts so they add up correctly using least squares optimization (i.e., find closest reconciled forecasts to the original forecasts).
- 3 This is available using `reconcile()`.

Forecast reconciliation

```
{r tourismets_reconciled, message=FALSE} tourism |> aggregate_key(Purpose * (State / Region), Trips = sum(Trips)) |> model(ets = ETS(Trips)) |> reconcile(ets_adjusted = min_trace(ets)) |> forecast(h = 2)
```

Hierarchical and grouped time series

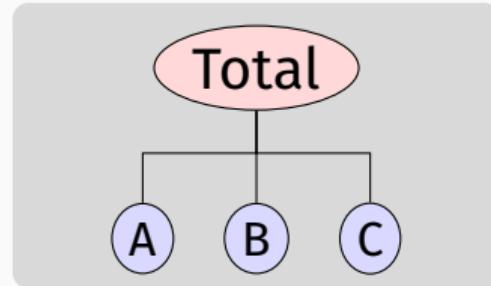
Every collection of time series with aggregation constraints can be written as

$$\mathbf{y}_t = \mathbf{S}\mathbf{b}_t$$

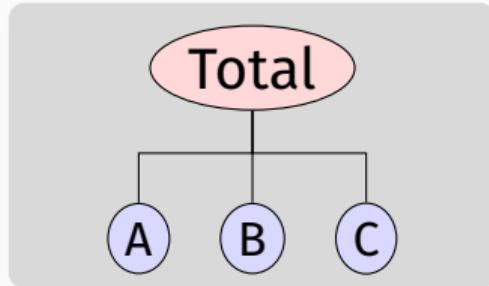
where

- \mathbf{y}_t is a vector of all series at time t
- \mathbf{b}_t is a vector of the most disaggregated series at time t
- \mathbf{S} is a “summing matrix” containing the aggregation constraints.

Hierarchical time series



Hierarchical time series

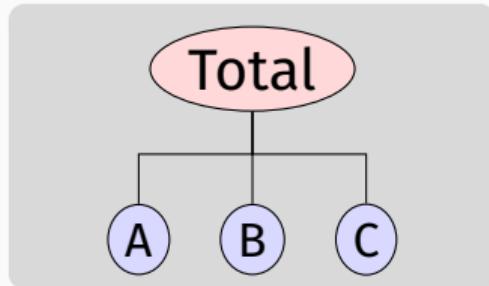


y_t : observed aggregate of all series at time t .

$y_{X,t}$: observation on series X at time t .

b_t : vector of all series at bottom level in time t .

Hierarchical time series



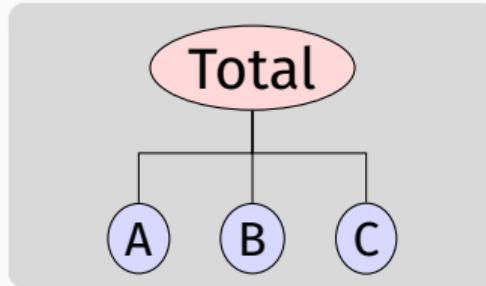
y_t : observed aggregate of all series at time t .

$y_{X,t}$: observation on series X at time t .

b_t : vector of all series at bottom level in time t .

$$\mathbf{y}_t = \begin{pmatrix} y_t \\ y_{A,t} \\ y_{B,t} \\ y_{C,t} \end{pmatrix} = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} y_{A,t} \\ y_{B,t} \\ y_{C,t} \end{pmatrix}$$

Hierarchical time series



y_t : observed aggregate of all series at time t .

$y_{X,t}$: observation on series X at time t .

\mathbf{b}_t : vector of all series at bottom level in time t .

$$\mathbf{y}_t = \begin{pmatrix} y_t \\ y_{A,t} \\ y_{B,t} \\ y_{C,t} \end{pmatrix} = \underbrace{\begin{pmatrix} 1 & 1 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}}_S \underbrace{\begin{pmatrix} y_{A,t} \\ y_{B,t} \\ y_{C,t} \end{pmatrix}}_{\mathbf{b}_t}$$

$$\mathbf{y}_t = S\mathbf{b}_t$$

Forecasting notation

Let $\hat{\mathbf{y}}_n(h)$ be vector of initial h -step forecasts, made at time n , stacked in same order as \mathbf{y}_t .

Forecasting notation

Let $\hat{\mathbf{y}}_n(h)$ be vector of initial h -step forecasts, made at time n , stacked in same order as \mathbf{y}_t .
(In general, they will not “add up’’.)

Forecasting notation

Let $\hat{\mathbf{y}}_n(h)$ be vector of initial h -step forecasts, made at time n , stacked in same order as \mathbf{y}_t .
(In general, they will not “add up’’.)

Reconciled forecasts must be of the form:

$$\tilde{\mathbf{y}}_n(h) = \mathbf{S}\mathbf{G}\hat{\mathbf{y}}_n(h)$$

for some matrix \mathbf{G} .

Forecasting notation

Let $\hat{\mathbf{y}}_n(h)$ be vector of initial h -step forecasts, made at time n , stacked in same order as \mathbf{y}_t .
(In general, they will not “add up”.)

Reconciled forecasts must be of the form:

$$\tilde{\mathbf{y}}_n(h) = \mathbf{S}\mathbf{G}\hat{\mathbf{y}}_n(h)$$

for some matrix \mathbf{G} .

- \mathbf{G} extracts and combines base forecasts $\hat{\mathbf{y}}_n(h)$ to get bottom-level forecasts.
- \mathbf{S} adds them up

Optimal combination forecasts

Main result

The best (minimum sum of variances) unbiased forecasts are obtained when $\mathbf{G} = (\mathbf{S}'\Sigma_h^{-1}\mathbf{S})^{-1}\mathbf{S}'\Sigma_h^{-1}$, where Σ_h is the h -step base forecast error covariance matrix.

Optimal combination forecasts

Main result

The best (minimum sum of variances) unbiased forecasts are obtained when $\mathbf{G} = (\mathbf{S}'\Sigma_h^{-1}\mathbf{S})^{-1}\mathbf{S}'\Sigma_h^{-1}$, where Σ_h is the h -step base forecast error covariance matrix.

$$\tilde{\mathbf{y}}_n(h) = \mathbf{S}(\mathbf{S}'\Sigma_h^{-1}\mathbf{S})^{-1}\mathbf{S}'\Sigma_h^{-1}\hat{\mathbf{y}}_n(h)$$

Problem: Σ_h hard to estimate, especially for $h > 1$.

Solutions:

- Ignore Σ_h (OLS) [`min_trace(method='ols')`]
- Assume $\Sigma_h = k_h\Sigma_1$ is diagonal (WLS)
[`min_trace(method='wls')`]
- Assume $\Sigma_h = k_h\Sigma_1$ and estimate it (GLS)
[`min_trace(method='shrink')` (the default)]

Features

- Covariates can be included in initial forecasts.
- Adjustments can be made to initial forecasts at any level.
- Very simple and flexible method. Can work with *any* hierarchical or grouped time series.
- Conceptually easy to implement: regression of base forecasts on structure matrix.

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Example: Australian tourism

```
tourism_agg <- tourism |>
  aggregate_key(Purpose * (State / Region),
    Trips = sum(Trips)
  )
fc <- tourism_agg |>
  filter_index(. ~ "2015 Q4") |>
  model(ets = ETS(Trips)) |>
  reconcile(ets_adjusted = min_trace(ets)) |>
  forecast(h = "2 years")
```

Example: Australian tourism

```
{r fctourism2, dependson='fctourism'} fc |> filter(is_aggregated(Purpose) &  
is_aggregated(State)) |> autoplot(tourism_agg, level = 95)
```

Example: Australian tourism

```
{r fctourism3, dependson='fctourism'} fc |> filter(is_aggregated(Purpose) &  
State == "VIC" & is_aggregated(Region)) |> autoplot(tourism_agg, level = 95)
```

Example: Australian tourism

```
{r fctourism4, dependson='fctourism'} fc |> filter(is_aggregated(Purpose) &  
Region == "Melbourne") |> autoplot(tourism_agg, level = 95)
```

Example: Australian tourism

```
{r fctourism5, dependson='fctourism'} fc |> filter(is_aggregated(Purpose) &  
Region == "Snowy Mountains") |> autoplot(tourism_agg, level = 95)
```

Example: Australian tourism

```
{r fctourism6, dependson='fctourism'} fc |> filter(Purpose == "Holiday" &  
Region == "Barossa") |> autoplot(tourism_agg, level = 95)
```

Example: Australian tourism

```
{r fctourism7, dependson='fctourism'} fc |> filter(is_aggregated(Purpose) &  
Region == "MacDonnell") |> autoplot(tourism_agg, level = 95)
```

Example: Australian tourism

```
fc <- tourism_agg |>
  filter_index(. ~ "2015 Q4") |>
  model(
    ets = ETS(Trips),
    arima = ARIMA(Trips)
  ) |>
  mutate(
    comb = (ets + arima) / 2
  ) |>
  reconcile(
    ets_adj = min_trace(ets),
    arima_adj = min_trace(arima),
    comb_adj = min_trace(comb)
  ) |>
  forecast(h = "2 years")
```

Forecast evaluation

```
{r fcaccuracy, dependson='fctourismcomb'} fc |> accuracy(tourism_agg)
```

Forecast evaluation

```
{r fcaccuracy2, dependson='fctourismcomb'} fc |>  
accuracy(tourism_agg) |> group_by(.model) |> summarise(MASE =  
mean(MASE)) |> arrange(MASE)
```

Outline

- 1 Time series data and tsibbles
- Example: Australian prison population
- Example: Australian pharmaceutical sales
- Lab Session 1
- Time plots
- Lab Session 2
- Seasonal plots
- Lab Session 3
- Seasonal or cyclic?
- Lag plots and autocorrelation
- Lab Session 4
- White noise
- Lab Session 5
- Per capita adjustments
- Lab Session 6
- Inflation adjustments
- Mathematical transformations
- Lab Session 7
- Time series decompositions
- Lab Session 8
- Multiple seasonality
- A cautionary tale
- Notice: Material planned to change
- STL features
- Lab Session 9
- Dimension reduction for features

Lab Session 20

- Prepare aggregations of the PBS data by Concession, Type, and ATC1.
- Use forecast reconciliation with the PBS data, using ETS, ARIMA and SNAIVE models, applied to all but the last 3 years of data.
- Which type of model works best?
- Does the reconciliation improve the forecast accuracy?
- Why doesn't the reconciliation make any difference to the SNAIVE forecasts?